CS3210 – Parallel Computing
Course Admin

Lecture 0

Overview

- Teaching Staff
- Course Overview:
  - Objectives
  - Topics & Lecture Schedule
- Course Structure:
  - Tutorial & Lab
  - Assessment
- References

Teaching Staff – **Lecturer**

- Dr. Cristina Carbunaru
- Contact Information:
  Email: **ccris@comp.nus.edu.sg**

  - Comments / suggestions welcome

  - Email in advance to book a consultation slot

Teaching Assistants

Name	Email
Sriram Sami	sriramsami@nus.edu.sg
Richard Willie	e0550368@u.nus.edu
Yap Zhi Heng	yap.zhh@u.nus.edu
Liew Zhao Wei	zhaoweiliew@gmail.com
Simon Julian Lauw	simonjulianl@u.nus.edu
Vimuth Mendis	e0550551@u.nus.edu
Theodore Leebrant	theodoreleebrant@u.nus.edu
Rohit Rajesh Bhat	rajeshbhatrohit@gmail.com
Seow Alex	seowalex@u.nus.edu
Koh Hong Po	kohhongpo@u.nus.edu

Course Objectives

- Provide an introduction to the field of parallel computing with **hands-on** parallel programming experience on **real parallel machines**

- Four major parts:
  1. Parallel architectures
  2. Parallel computation models
  3. Parallel algorithm design and programming
  4. Performance of parallel programs

Not-so-formal Objectives

- Designing and writing scalable parallel programs
  - Parallel thinking and writing code in popular parallel programming languages
  - Emphasis on performance
- Parallel computer hardware implementation: how parallel computers work
  - Characteristics of the machine really matter
- Efficiency of running parallel programs on parallel hardware
  - Efficient != fast: running faster on a parallel computer, it does not mean it is using the hardware efficiently

Topics in Lectures

- Background on parallelism (from OS)
  - L2: Processes and threads
- Architecture
  - L3: Processor and memory organization
  - L7: Cache coherence and memory consistency
  - L11: Interconnection networks
- Parallel computation models
  - L4: Shared-memory programming models
  - L6: Data parallel models (GPGPU)
  - L9, L10: Distributed-memory programming models
- Performance and scalability of parallel programs
  - L5: Performance of parallel systems
  - L8: Performance instrumentation
- New trends
  - L12: Energy efficient computing

Overview

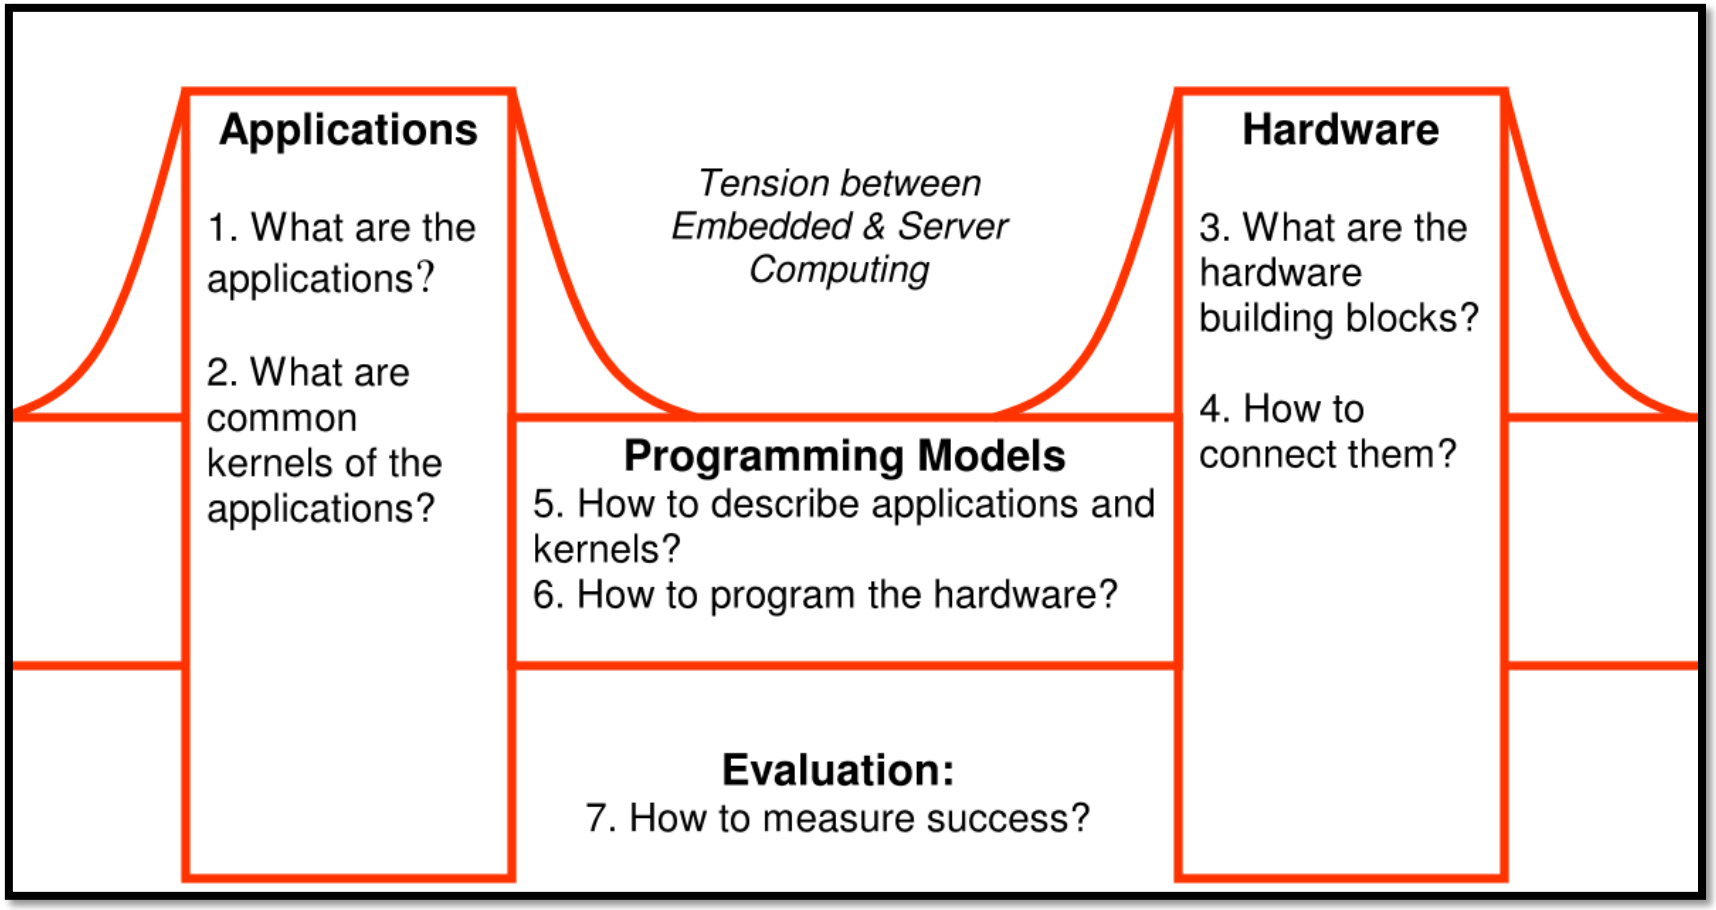

Picture taken from "The Landscape of Parallel Computing Research: A View from Berkeley" (2008)

Program Parallelization: Steps

- 3 main steps:

1. **Decomposition** of the computations
2. **Scheduling** (assignment of tasks to processes (or threads))
3. **Mapping** of processes (or threads) to physical processors (or cores)

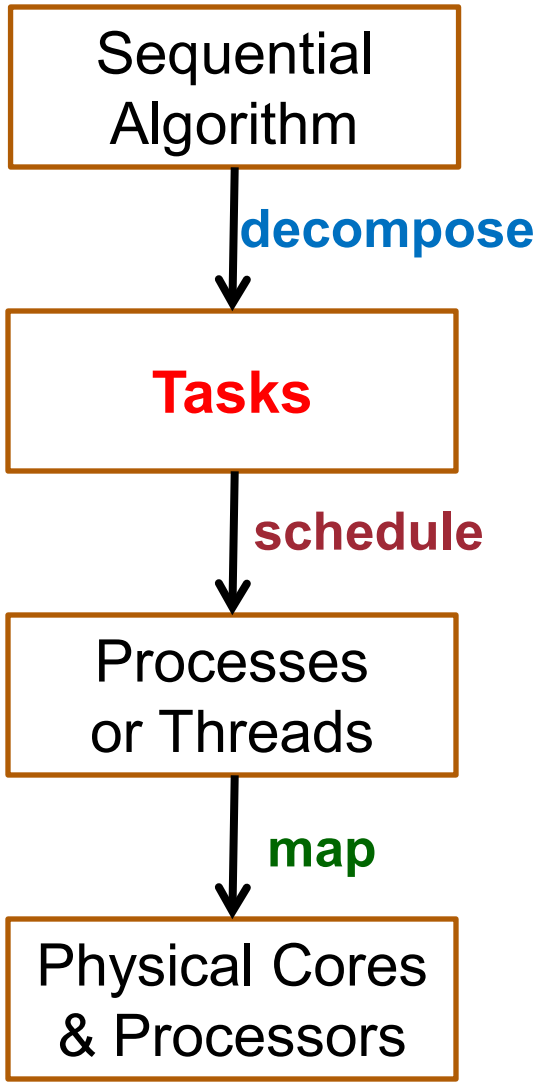

High-level Structure of the Module

- L1-L7: **Shared-memory models**
  - Architecture, memory consistency, programming
  - OpenMP and CUDA
- L8-L11: **Distributed memory models**
  - Architecture, interconnects, programming
  - MPI
- with an emphasis on **Parallel performance**

In-person Teaching

- All lectures are **recorded**
  - Recordings published in Canvas
  - No attendance taking at lecture
- Tutorial and labs: in-person, in various tutorial rooms
  - Attendance is taken
  - Lab sheets submissions

Course Structure – Tutorial & Lab

- Tutorials and Labs in different weeks:
  - Starting with Lab 1 in Week 3, Tutorial 1 in Week 4, …., etc
    - Register on EduRec for a tutorial slot
  - Each session is 2 hours long
    - 2nd hour is open session
  - Parallel & Distributed Computing Lab, Com 1, #B1-02
    - Classes are conducted in other tutorial rooms

Activities - Assessment

- Continuous Assessment (60%)
  - Lab (6%) – 3 lab sheets
  - Tutorial (4%) – attendance and participation
  - Programming Assignments (35%)
    - C/C++
    - Work in pairs or individually
    - Teammates might be registered to different tutorial slots
  - 4 Quizzes (15%)
    - Weeks 5, 8, 10,12
    - **Due on Sun, 8pm**
- Exam (40%) – Wed, 29 Nov, 9am
  - Open book, **in-person**

Schedule

Week	Lecture (Mon, 2pm) – SR1 (COM1-206)	Tutorial (Mon-Thu)	Lab (Mon-Thu)	Deadlines
1	L0: Course Admin L1: Introduction			
2	L2: Processes, Thread & Synchronization			
3	L3: Processor & Memory Organization		L1: Processes & Threads	
4	L4: Parallel Programming Models – I	T1: Parallel Computer Architecture		Wed, 2pm – Lab 1 due
5	L5: Performance of Parallel Systems		L2: Shared-memory Programming, Performance Instrumentation	Sun, 8pm – Quiz 1 due
6	L6: GPU Programming	T2: Performance of Parallel Systems		Wed, 2pm – Lab 2 due
Recess				Wed, 2pm – Assignment 1
7	L7: Cache Coherence & Memory Consistency		L3: CUDA programming	
8	L8: Performance Instrumentation	T3: Memory Consistency & CUDA Programming		Sun, 8pm – Quiz 2 due
9	L9: Parallel Programming Models – II	T4: Shared-memory Problems		
10	L10: Message-passing Programming		L4: Introduction to Distributed-memory Programming	Mon, 2pm - Assignment 2 Sun, 8pm – Quiz 3 due
11	L11: Interconnection Networks		L5: Message-passing Programming with MPI	Wed, 2pm – Lab 4 due
12	L12: Energy-efficient Computing and Cloud Computing L13: Summary & Recap	T5: MPI		Sun, 8pm – Quiz 4 due
13	PH: No lecture		L6: Benchmarks & performance	Fri, 2pm – Assignment 3 (optional) Sun, 2pm – Lab 6 due
Exam	29 Nov, 9am - Exam In-person, open book			Wed, 29 Nov, 9am - Exam

Reference book

- Robert Robey, Yuliana Zamora, ”Parallel and High-Performance Computing”, Manning Publications

Supplementary Reference Book

- Parallel Programming for Multicore and Cluster Systems
  - by Thomas Rauber and Gudula Rünger
  - 1st Edition, Springer-Verlag, 2010
  - Ebook - NUS Digital Library - http://linc.nus.edu.sg/record=b2974382

- 2nd Edition (2013) available online