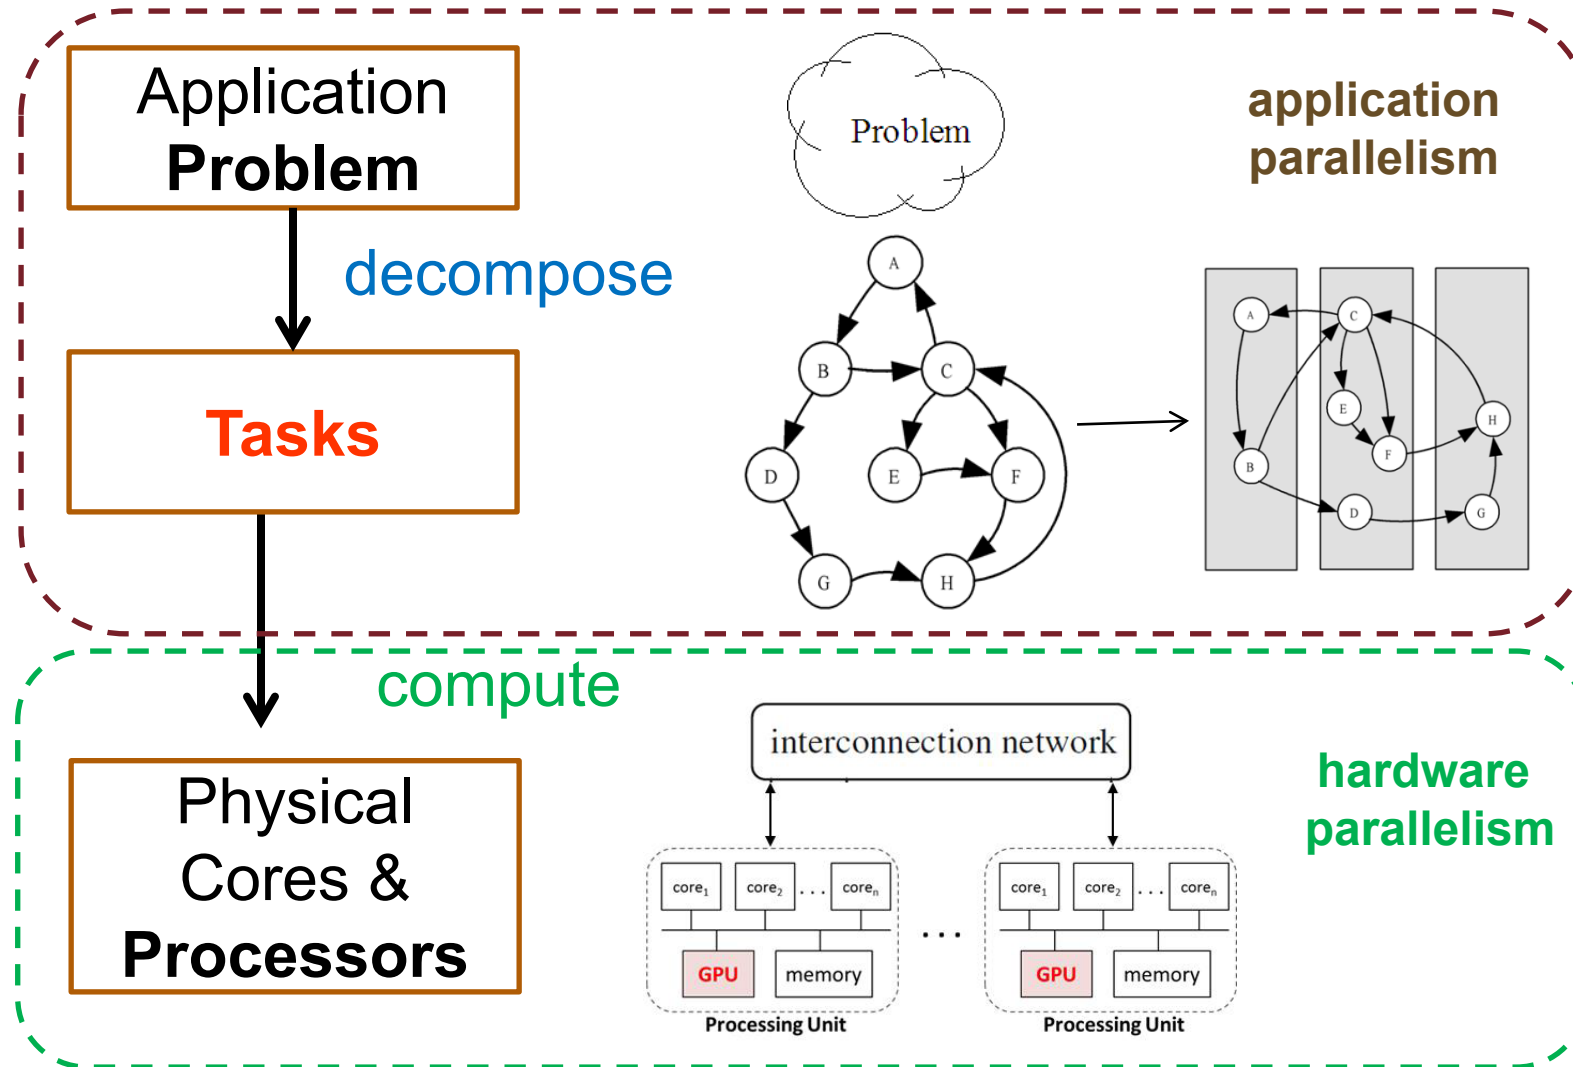


Parallel Computing Architectures

Lecture 03

Parallel Computing



Computer architecture

- Key concepts about how modern computers work
 - Concerns on parallel execution
 - Challenges of accessing memory
- Understanding these architecture basics help you
 - **Understand and optimize the performance** of your parallel programs
 - **Gain intuition** about what workloads might benefit from fast parallel machines

Outline

- Processor Architecture and Technology Trends
 - Various forms of parallelism
- Flynn's Parallel Architecture Taxonomy
- Architecture of Multicore Processors
- Memory Organization
 - Distributed-memory Systems
 - Shared-memory Systems
 - Hybrid (Distributed-Shared Memory) Systems

Concurrency vs. Parallelism

Concurrency

- Two or more tasks can start, run, and complete in overlapping time periods
- They might not be running (executing on CPU) at the same instant
- Two or more execution flows make **progress** at the same time by interleaving their executions or by executing instructions (on CPU) at exactly the same time

Parallelism


- Two or more tasks can run (execute) simultaneously, at the exact same time
- Tasks do not only make progress, they actually execute **simultaneously**

Source of Processor Performance Gain

- Parallelism of various forms are the main source of performance gain

- Let us understand parallelism at the:

- ❑ Bit Level
- ❑ Instruction Level
- ❑ Thread Level
- ❑ Processor Level:
 - Shared Memory
 - Distributed Memory



Bit Level Parallelism


- Word size may mean:
 - ❑ Unit of transfer between processor \leftrightarrow memory
 - ❑ Memory address space capacity
 - ❑ Integer size
 - ❑ Single precision floating point number size
- Word size trend:
 - ❑ Varied in the 50s to 70s
 - ❑ Following x86 processors:
 - 16 bits (8086 – 1978)
 - 32 bits (80386 – 1985)
 - 64 bits (Pentium 4 / Opteron – 2003)

Instruction Level Parallelism


- Execute instructions in parallel:
 - Pipelining (parallelism across time)
 - Superscalar (parallelism across space)
- Pipelining:
 - Split instruction execution in multiple stages, e.g.
 - Fetch (IF), Decode (ID), Execute (EX), Write-Back (WB)
 - Allow multiple instructions to occupy different stages in the same clock cycle
 - Provided there is no data / control dependencies
 - Number of pipeline stages == Maximum achievable speedup

Pipelined Execution: Illustration

Non-pipelined




Pipelined



- Disadvantages
 - ❑ Independence
 - ❑ Bubbles
 - ❑ Hazards: data and control flow
- Speculation
- Out-of-order execution
 - ❑ Read-after-write

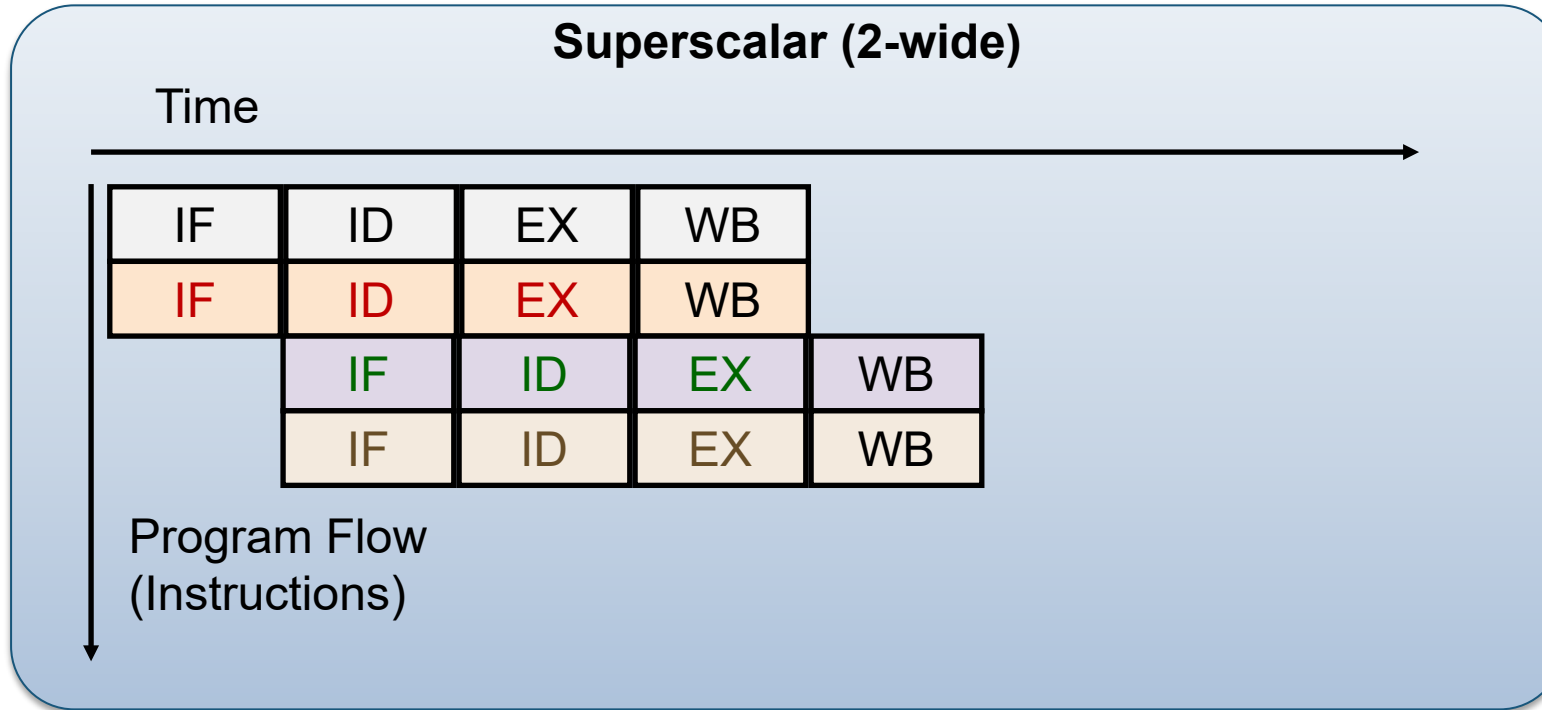
The end of ILP



Instruction Level Parallelism: **Superscalar**

- Duplicate the pipelines:
 - Allow multiple instructions to pass through the same stage
 - Scheduling is challenging (decide which instructions can be executed together):
 - Dynamic (Hardware decision)
 - Static (Compiler decision)
- Most modern processors are superscalar
 - e.g. each intel i7 core has 14 pipeline stages and can execute 6 micro-ops in the same cycle

Superscalar Execution: Illustration

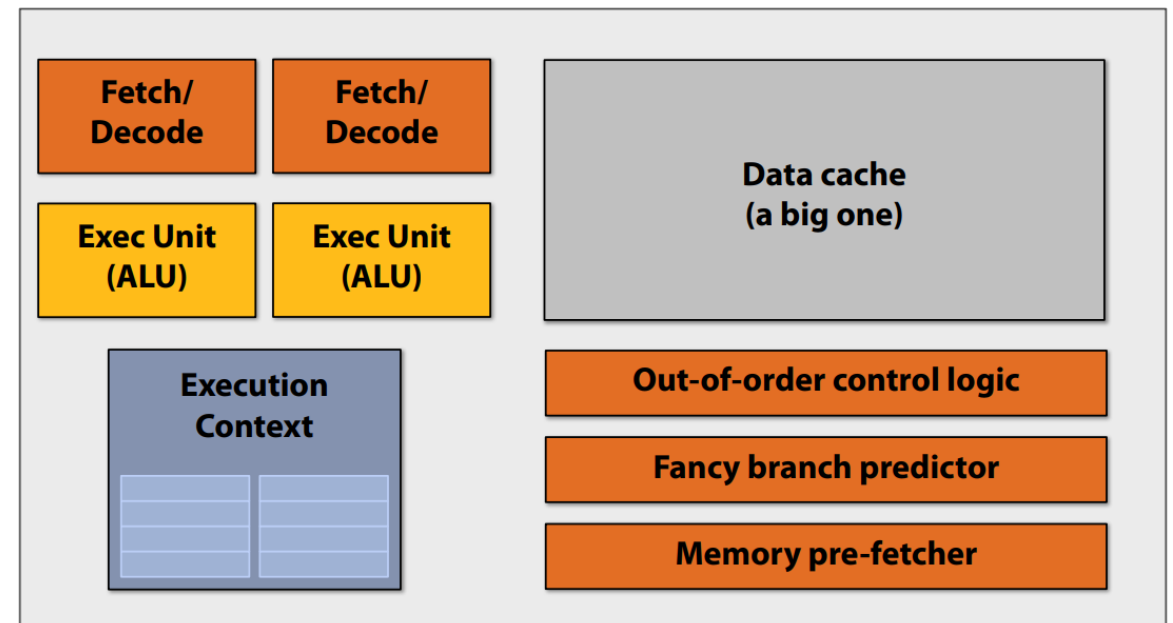
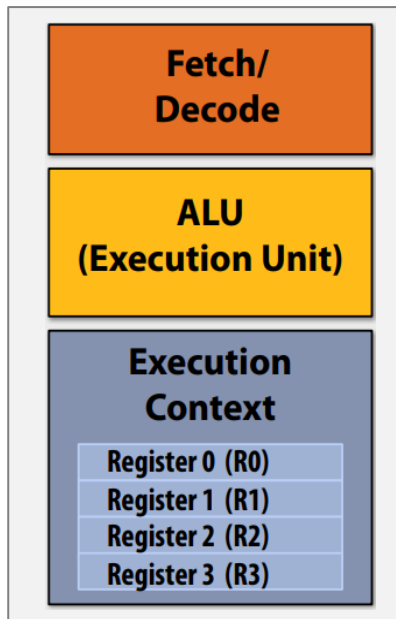


Disadvantages:
structural hazard

- Cycles-per-instruction (CPI) → Instructions-per-cycle (IPC)

Pipelined vs Superscalar Processor

- Determine what instruction to run next
 - Execution unit: performs the operation described by an instruction
 - Registers: store value of variables used as inputs and outputs to operations
- Processor automatically finds independent instructions in an instruction sequence and can execute them in parallel on execution units
 - **Instructions come from the same execution flow (thread)**



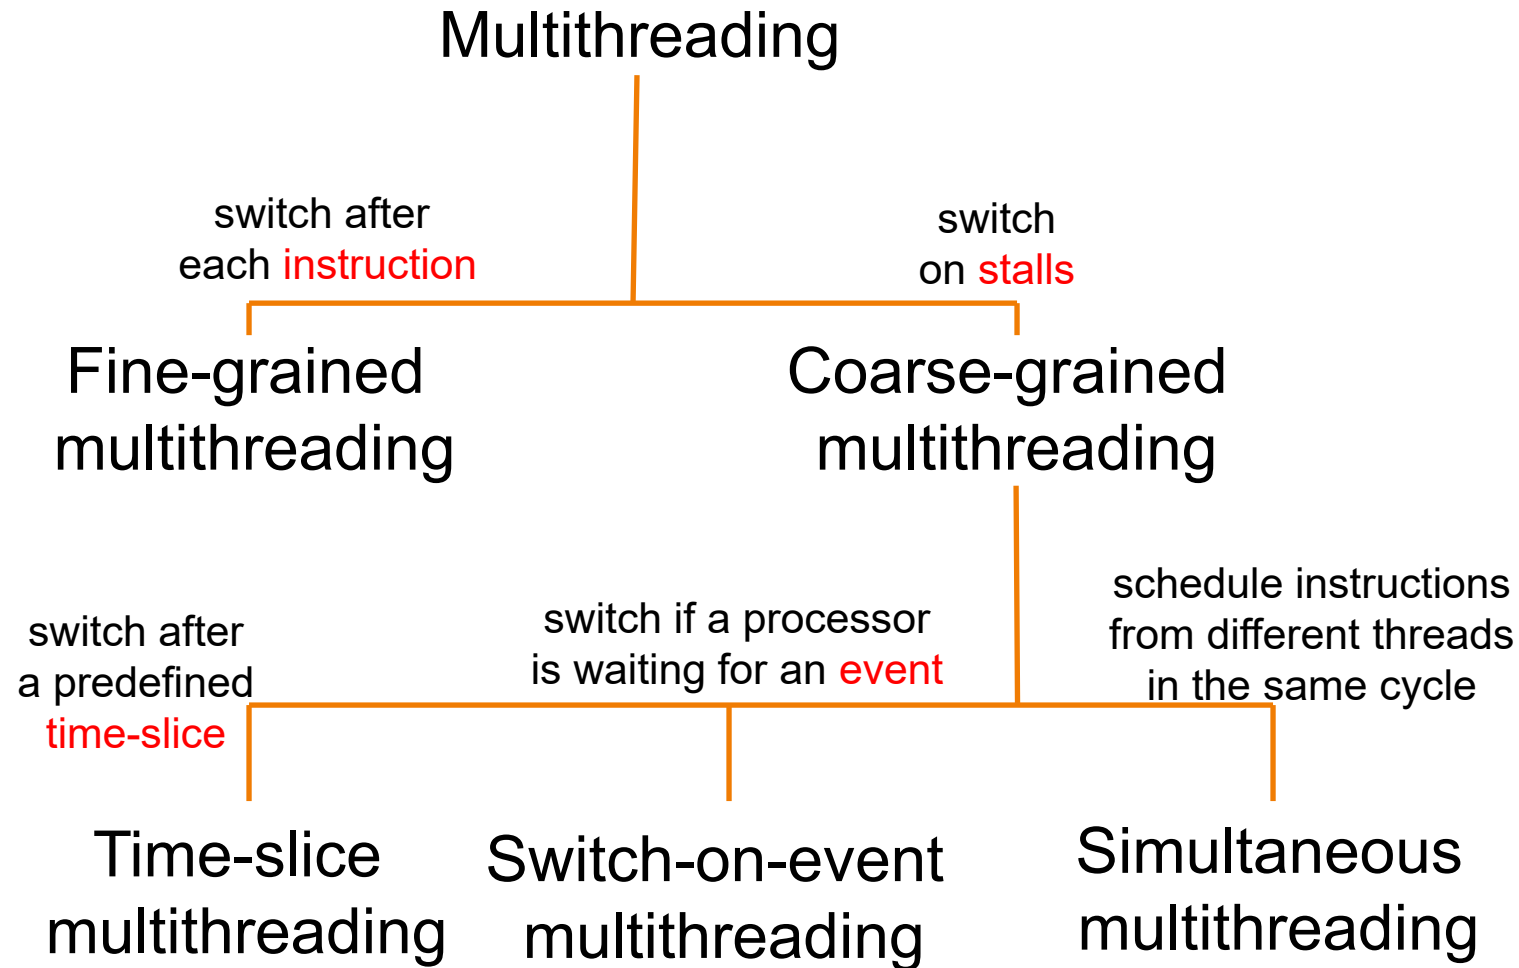
Thread Level Parallelism: Motivation

- Instruction-level parallelism is **limited**
 - ❑ For typical programs only 2-3 instructions can be executed in parallel (either pipelined / superscalar)
 - ❑ Due to data/control dependencies
- Multithreading was originally a software mechanism
 - ❑ Allow multiple parts of the same program to execute concurrently
- Key idea:
 - ❑ **The processor can execute the threads in parallel**

Thread Level Parallelism in the Processor

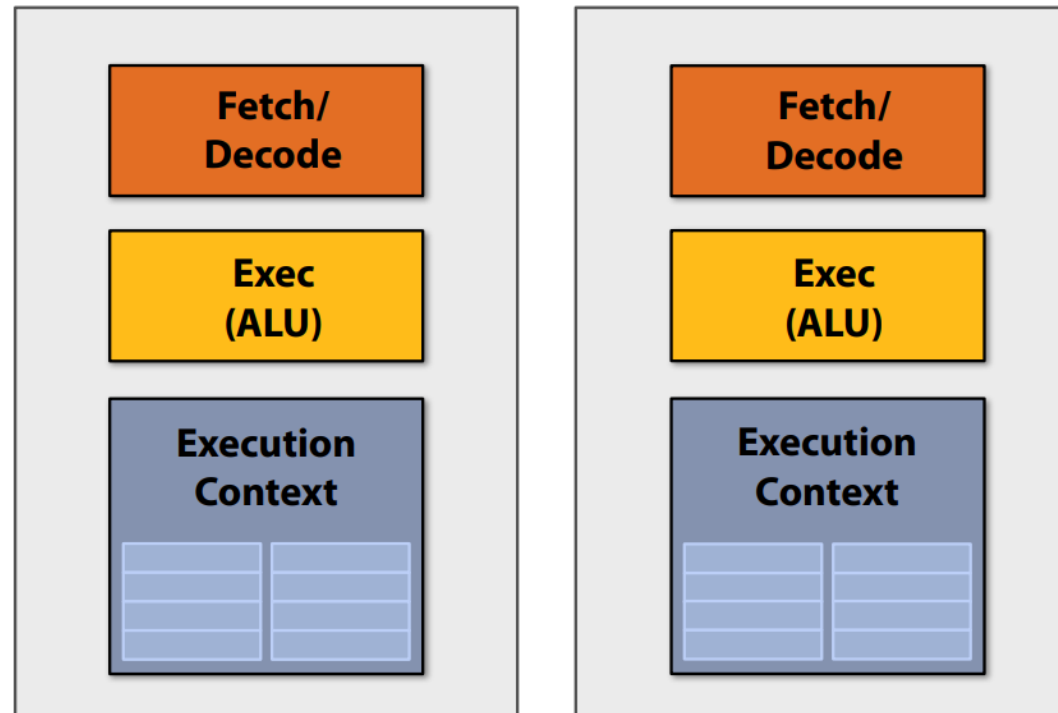
- Processor can provide hardware support for multiple "thread contexts"
 - Known as **simultaneous multithreading (SMT)**
 - Information specific to each thread, e.g. Program Counter, Registers, etc
 - Software threads can then execute in parallel
 - Many implementation approaches
- Example:
 - Intel processors with **hyper-threading** technology, e.g. each i7 core can execute 2 threads at the same time

Multithreading Implementations



Processor Level Parallelism (Multiprocessing)

- Add more cores to the processor
- The application should have multiple execution flows
 - ❑ Each process/thread needs an independent context that can be mapped to multiple processor cores




Flynn's Parallel Architecture Taxonomy

- One commonly used taxonomy of parallel architecture:
 - ❑ Based on the parallelism of instructions and data streams in the most constrained component of the processor
 - ❑ Proposed by M.Flynn in 1972(!)
- Instruction stream:
 - ❑ A single execution flow
 - ❑ i.e. a single Program Counter (PC)
- Data stream:
 - ❑ Data being manipulated by the instruction stream

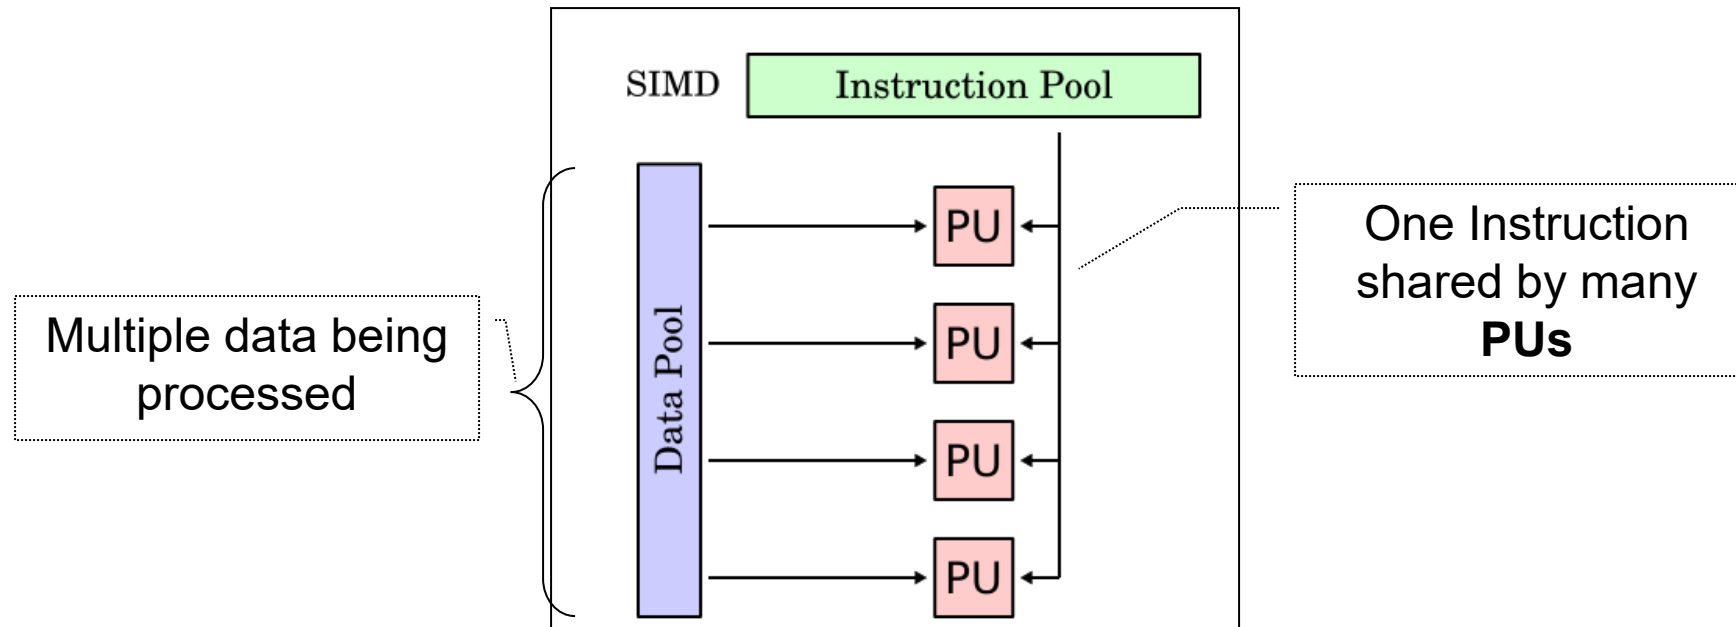
Single Instruction Single Data (SISD)

- A single instruction stream is executed
- Each instruction work on single data
- Most of the uniprocessors fall into this category



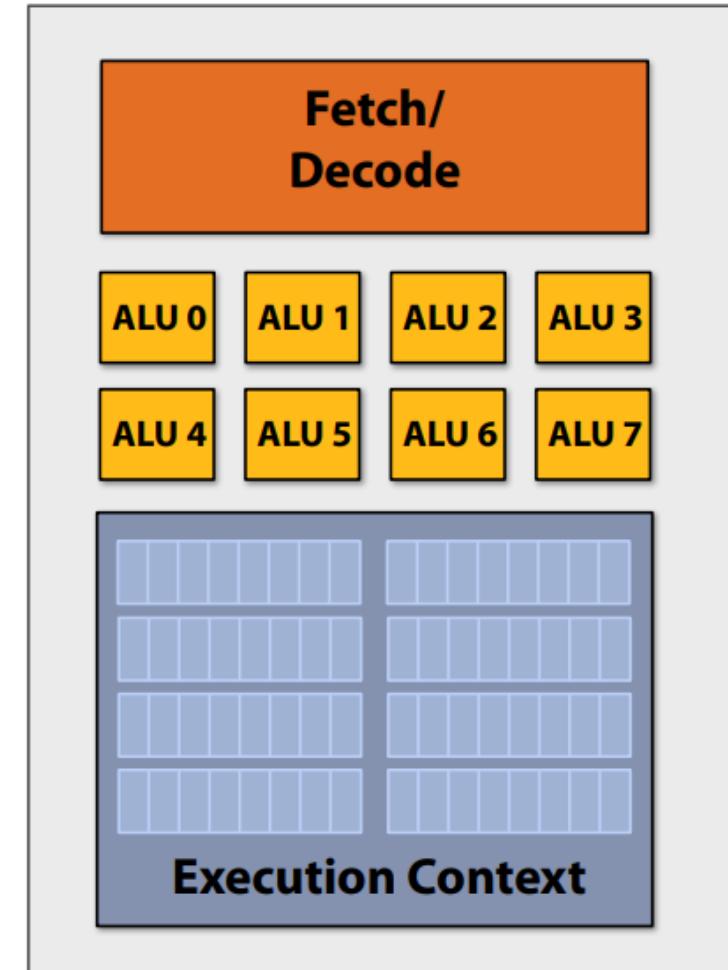
Single Instruction Multiple Data (SIMD)

- A single stream of instructions
- Each instruction works on multiple data
- Popular model for supercomputer during 1980s:
 - ❑ Exploit **data parallelism**, commonly known as vector processor
- Modern processor has some forms of SIMD:
 - ❑ E.g. the SSE, AVX instructions in intel x86 processors



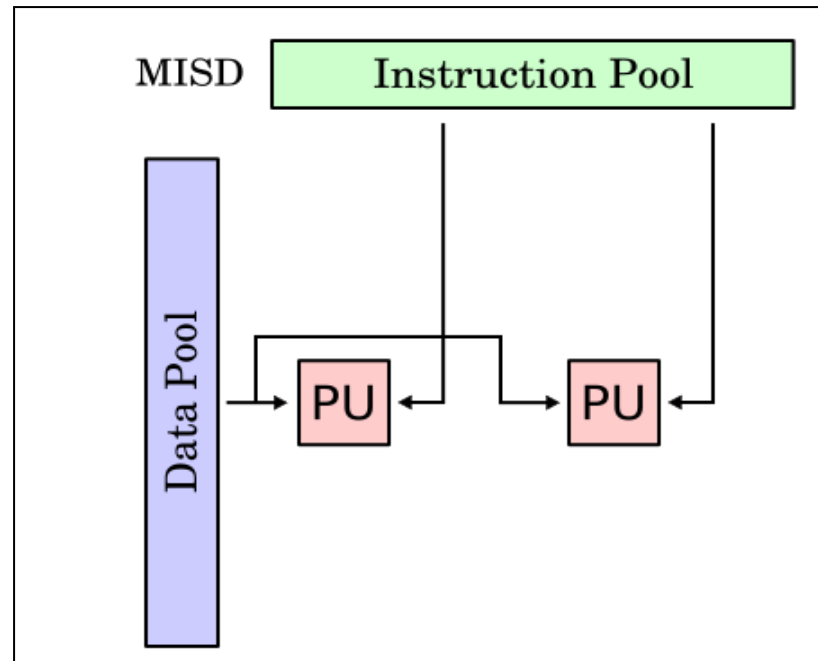
SIMD nowadays

- Data parallel architectures
 - ❑ AVX instructions
 - ❑ GPGPUs
- Same instruction broadcasted to all ALUs
- AVX: Intrinsic functions operate on vectors of four 64-bit values (e.g., vector of 4 doubles)
- Not great for divergent executions




Multiple Instruction Single Data (MISD)

- Multiple instruction streams
- All instruction work on the same data at any time
- No actual implementation except for the systolic array



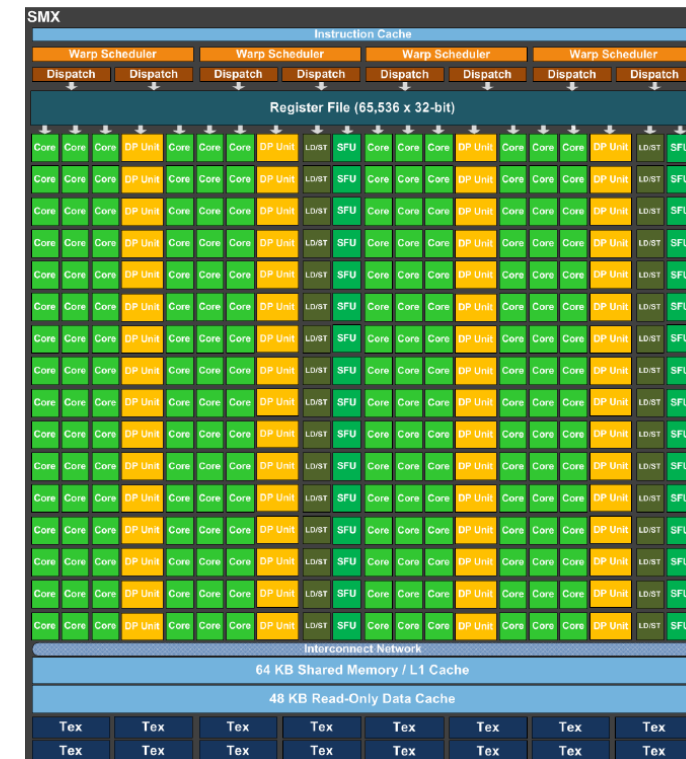
Multiple Instruction Multiple Data (MIMD)

- Each PU fetch its own instruction
- Each PU operates on its data
- Currently the most popular model for multiprocessor



Variant – SIMD + MIMD

- Stream processor (nVidia GPUs)
 - ❑ A set of threads executing the same code (effectively SIMD)
 - ❑ Multiple set of threads executing in parallel (effectively MIMD at this level)




MULTICORE ARCHITECTURE

Architecture of Multicore Processors

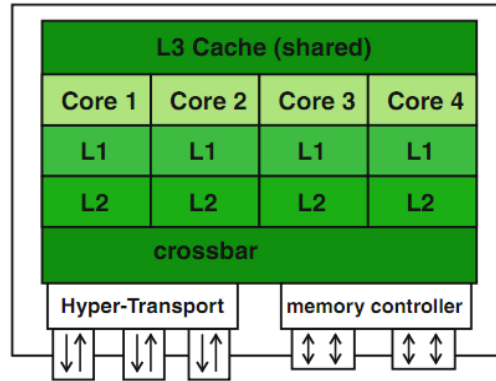
- Hierarchical design
- Pipelined design
- Network-based design

Hierarchical Design

- Multiple cores share multiple caches
- Cache size increases from the leaves to the root
- Each core can have a separate L1 cache and shares the L2 cache with other cores
- All cores share the common external memory
- Usages
 - ❑ Standard desktop
 - ❑ Server processors
 - ❑ Graphics processing units




Hierarchical Design - Examples

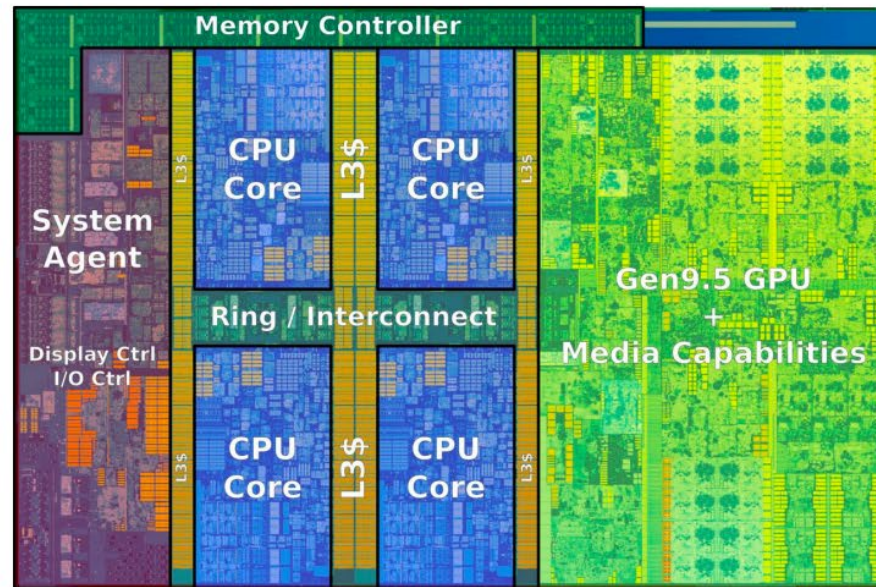


Quad-Core AMD Opteron


Each core is sophisticated, out-of-order processor to maximize ILP



Intel Quad-Core Xeon




Example: Intel Core i7 processor (Kaby Lake)



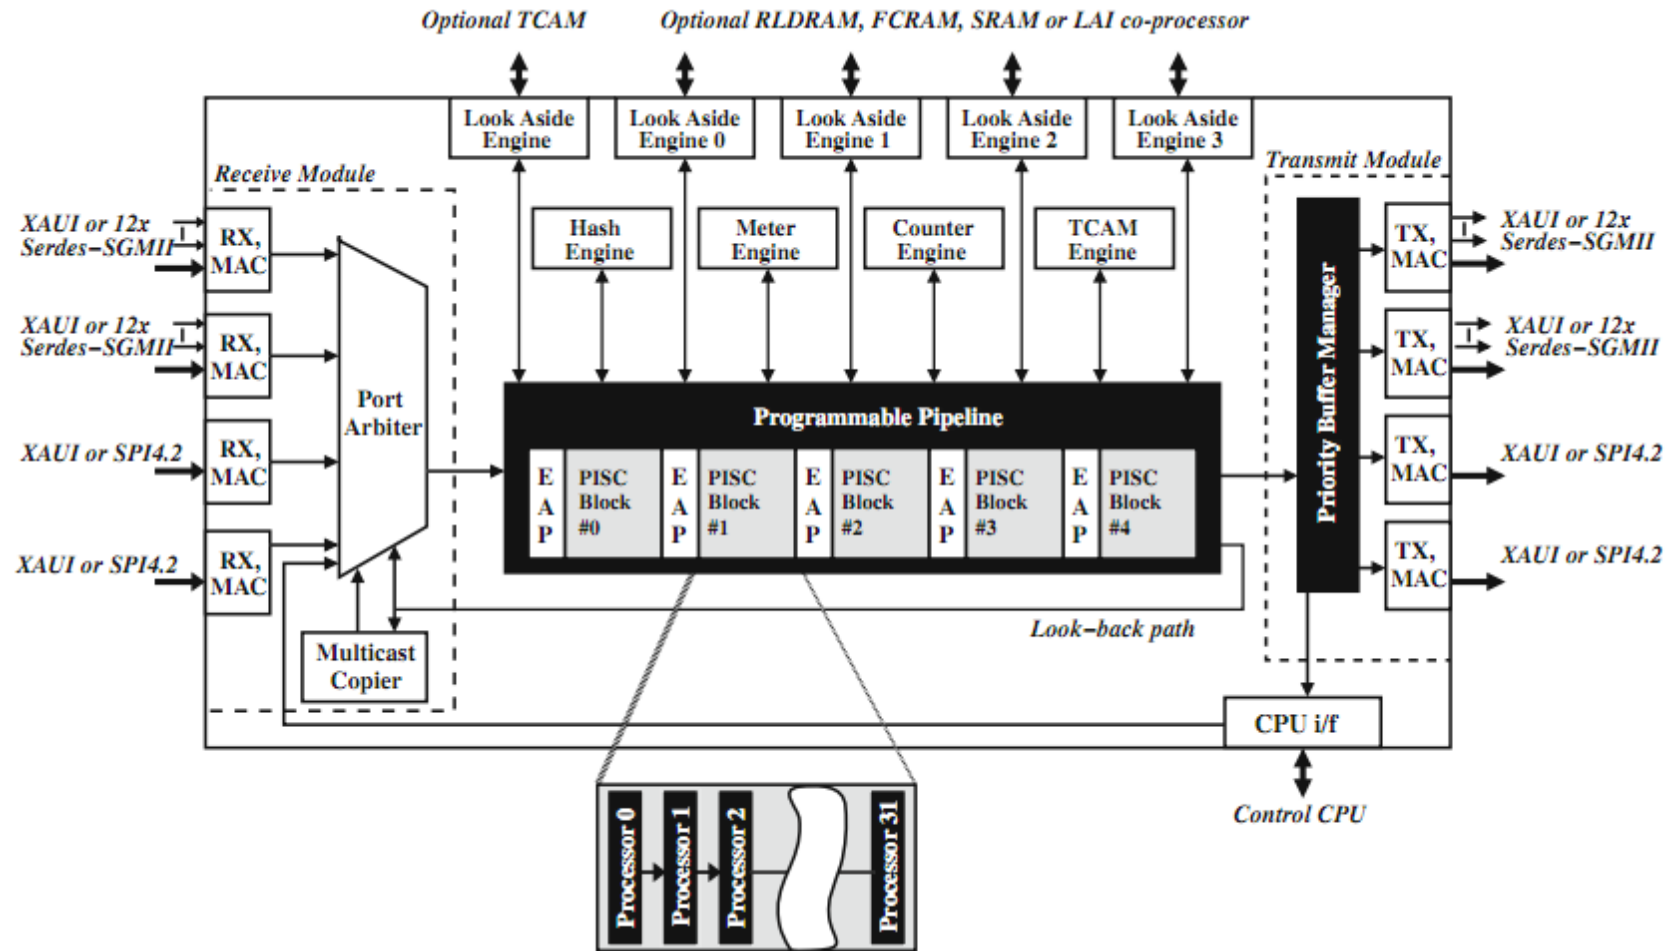
**Intel Core i7 (quad core)
(interconnect is a ring)**

Pipelined Design

- Data elements are processed by multiple execution cores in a **pipelined way**
- Useful if same computation steps have to be applied to a long sequence of data elements
 - ❑ E.g. processors used in routers and graphics processors



Example: Pipelined Design

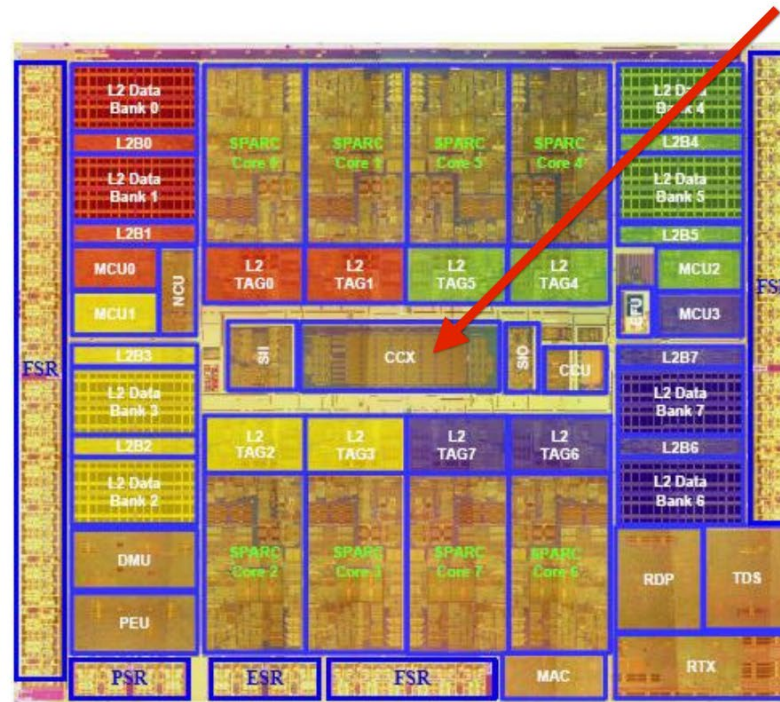



Xelerator X11 network processor

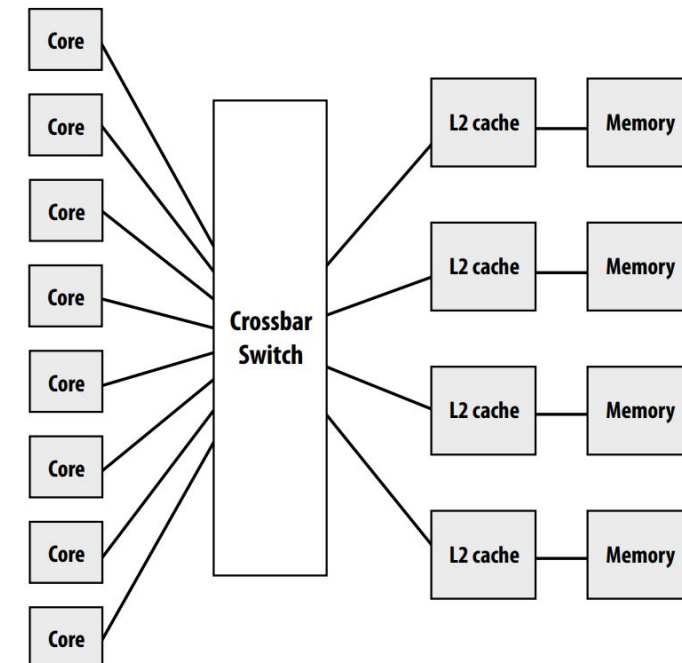
Network-Based Design

- Cores and their local caches and memories are connected via an interconnection network

SUN Niagara 2 (UltraSPARC T2)



Note area of crossbar (CCX):
about same area as one core on chip




Future Trends

- Efficient on-chip interconnection
 - ❑ Enough bandwidth for data transfers between the cores
 - ❑ Scalable
 - ❑ Robust to tolerate failures
 - ❑ Efficient energy management
 - ❑ Reduce memory access time
- Key word: Network on Chip (NoC)

MEMORY ORGANIZATION

Parallel Computer Component




Uniprocessor

- Typical uniprocessor components:
 - ❑ Processor
 - ❑ One or more level of caches
 - ❑ Memory module
 - ❑ Other (e.g. I/O)
- These components similarly present in a parallel computer setup
- Processors in a parallel computer systems is also commonly known as **processing element**

Recap: why do modern processors have cache?

- Processors run efficiently when data is resident in caches
 - Caches reduce memory access latency *

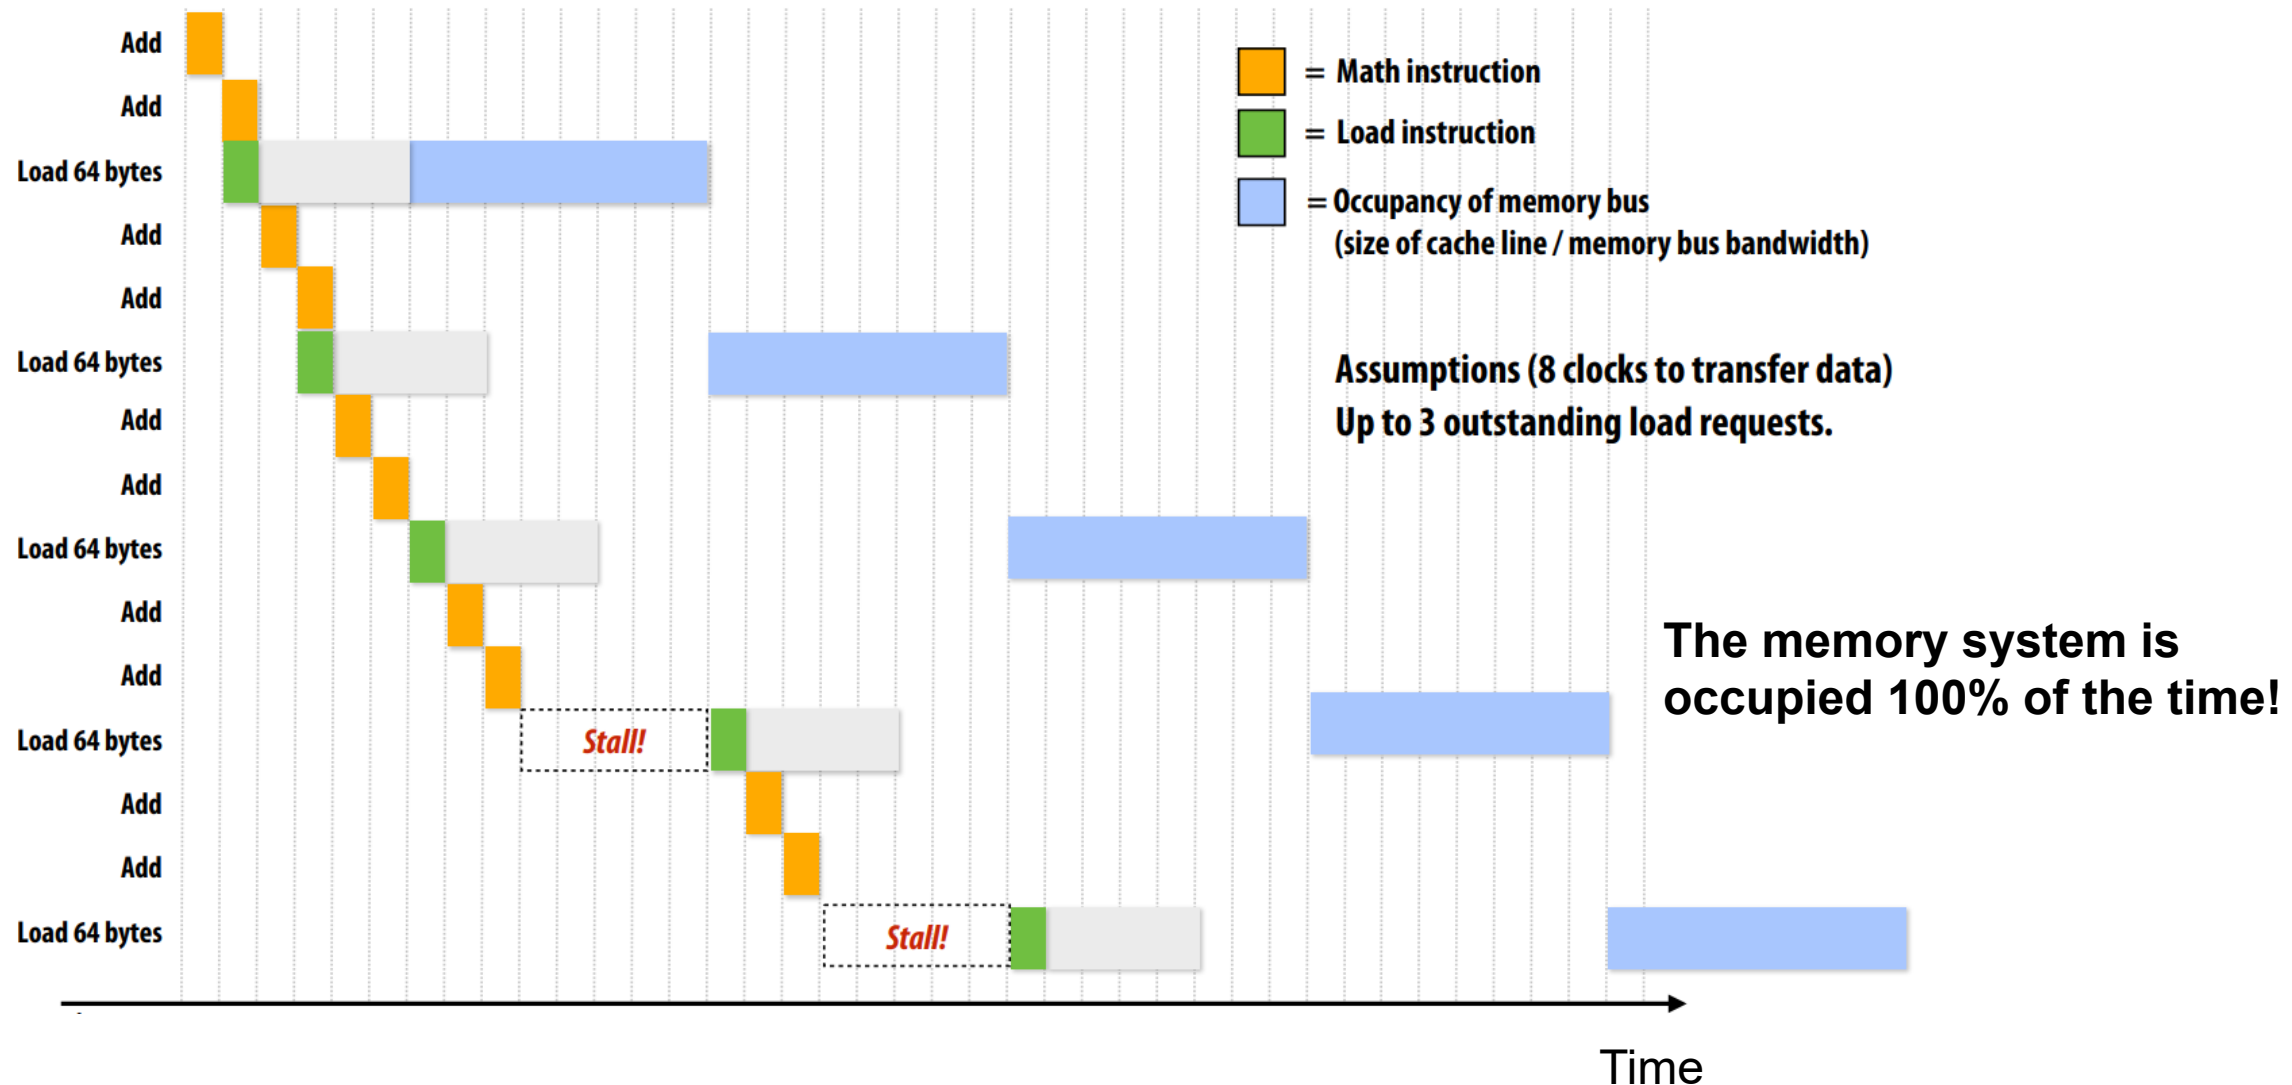
* Caches provide high bandwidth data transfer to CPU



Recap: memory latency and bandwidth

- Memory latency: the amount of time for a memory request (e.g., load, store) from a processor to be serviced by the memory system
 - Example: 100 cycles, 100 nsec
- Memory bandwidth: the rate at which the memory system can provide data to a processor
 - Example: 20 GB/s
- Processor “stalls” when it cannot run the next instruction in an instruction stream because of a dependency on a previous instruction

Execution on a Processor (one add per clock)

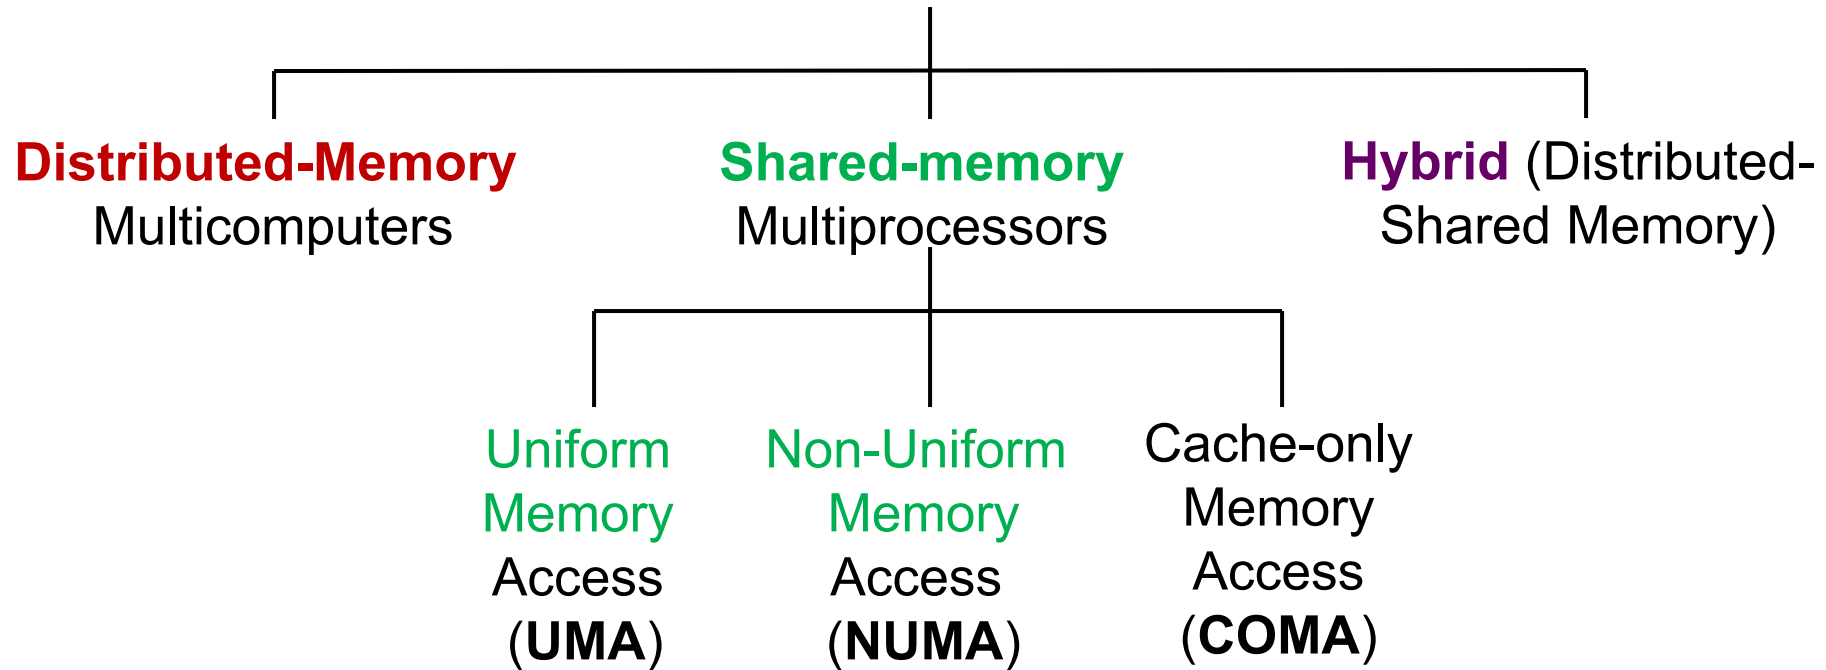


In Modern Computing, Bandwidth is the Critical Resource


- Performant parallel programs should:
 - Organize computation to fetch data from memory less often
 - Reuse data previously loaded by the same thread (temporal locality optimizations)
 - Share data across threads (inter-thread cooperation)
 - Favor performing additional arithmetic to storing/reloading values (the math is “free”)
 - **Main point:** programs must access memory infrequently to utilize modern processors efficiently

Memory Organization of Parallel Computers

Parallel Computers

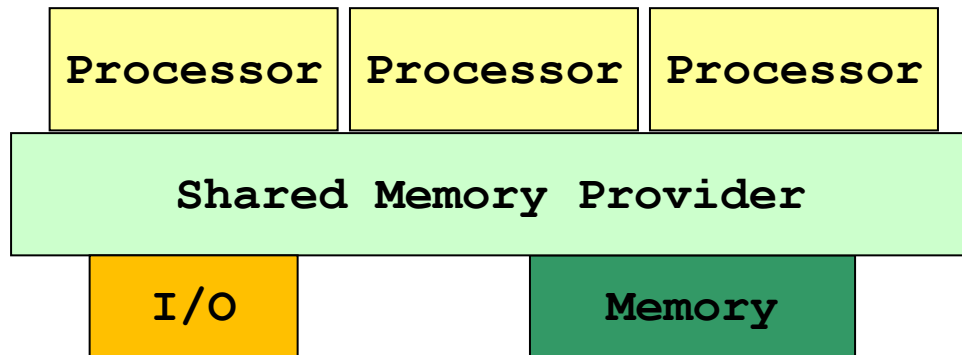


Distributed-Memory Systems




- Each node is an independent unit:
 - With processor, memory and, sometimes, peripheral elements
- Physically distributed memory module:
 - ➔ Memory in a node is private

Shared Memory System




- Parallel programs / threads access memory through the shared memory provider
- Program is unaware of the actual hardware memory architecture
 - Cache coherence and memory consistency

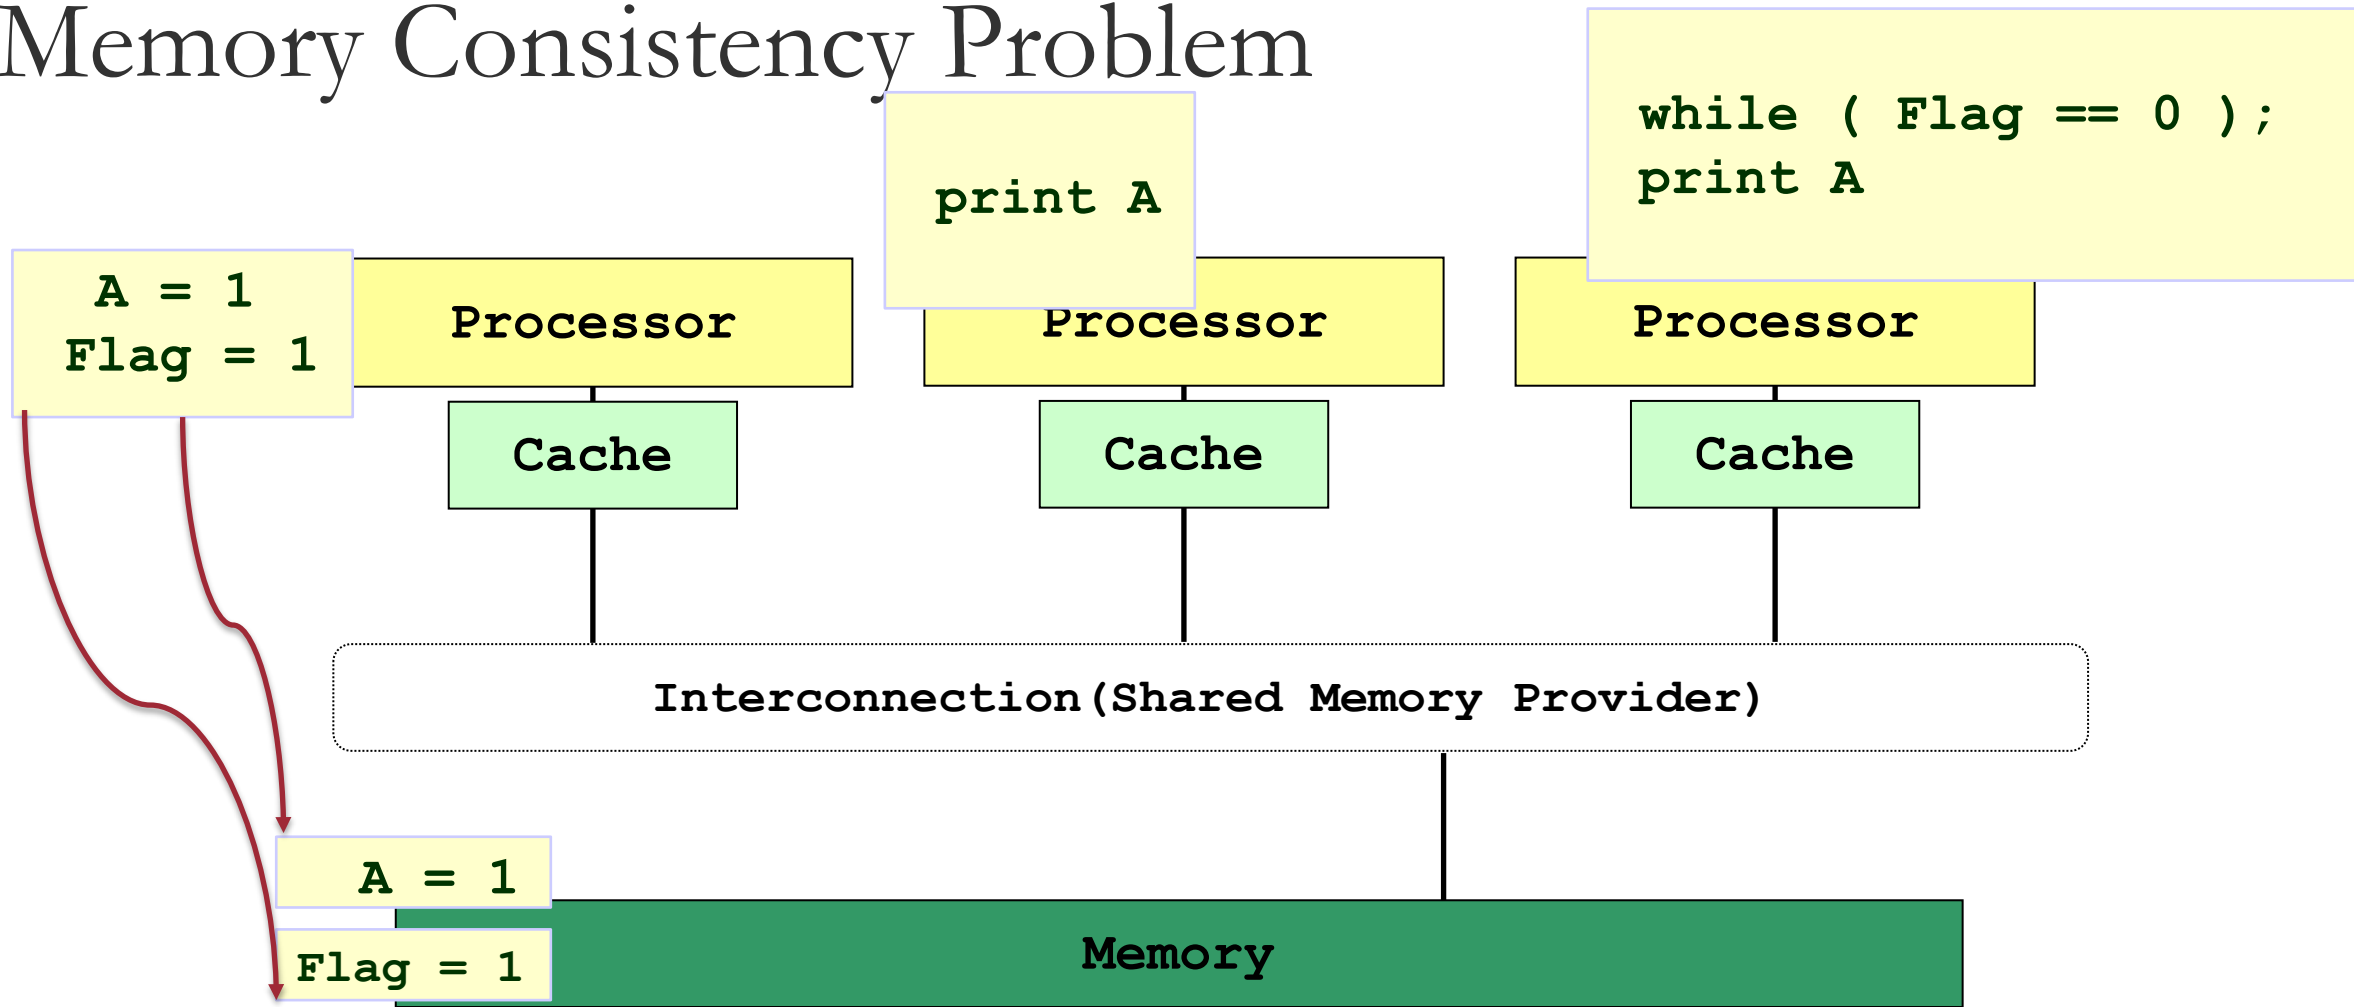


Cache Coherence

- Multiple copies of the same data exist on different caches
- Local update by processor → Other processors should not see the unchanged data



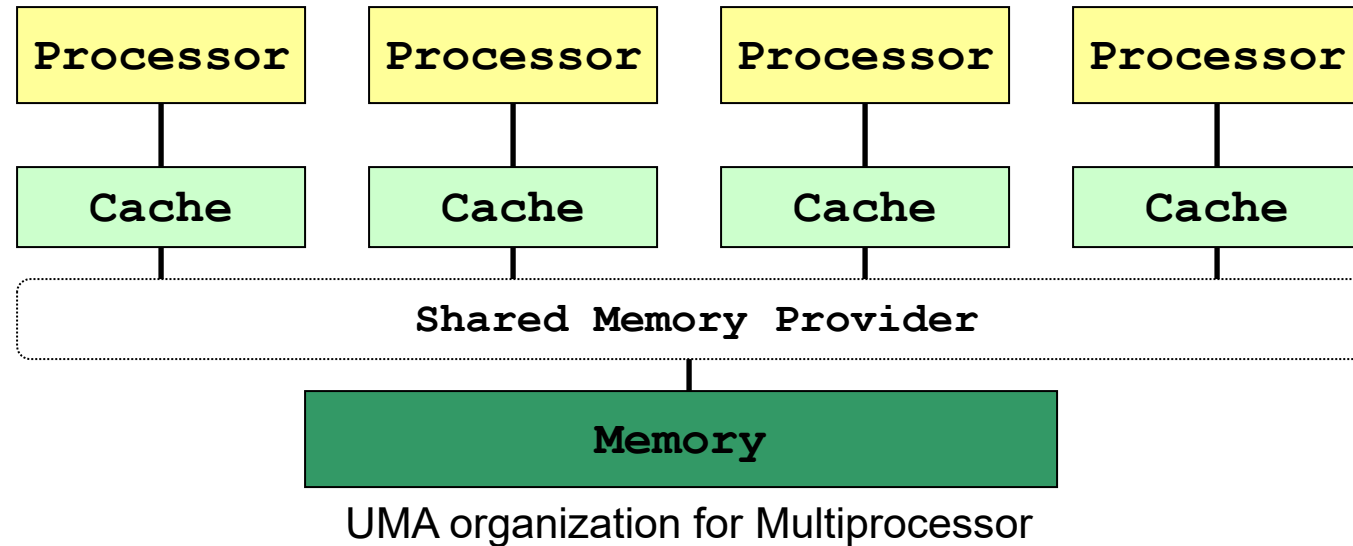
Memory Consistency Problem



Further Classification – Shared Memory

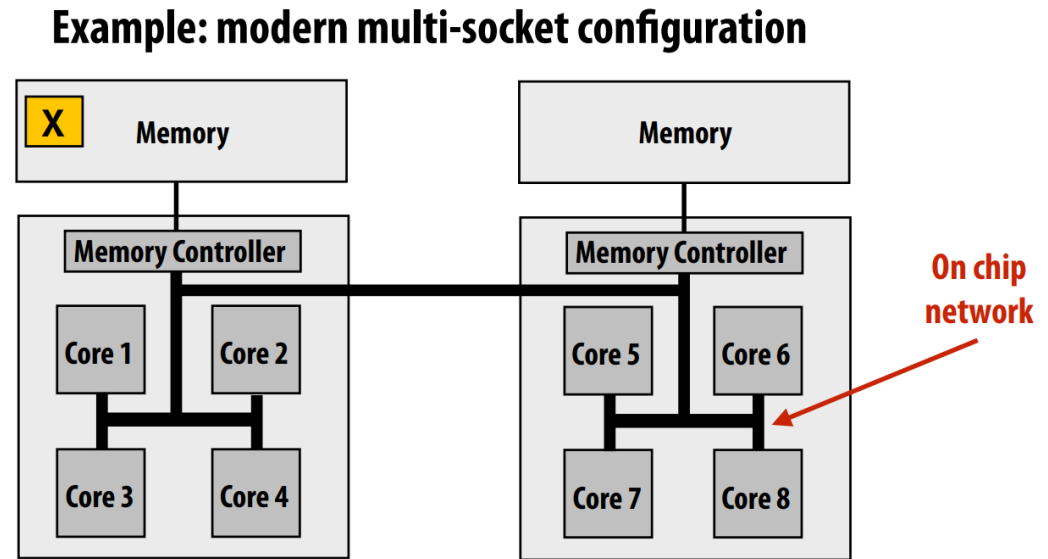

- Two factors can further differentiate **shared memory systems**:
 - Processor to Memory Delay (**UMA** / **NUMA**)
 - Whether delay to memory is uniform
 - Presence of a local cache with cache coherence protocol (**CC/NCC**):
 - Same shared variable may exist in multiple caches
 - Hardware ensures correctness via cache coherence protocol

Uniform Memory Access (Time) (UMA)



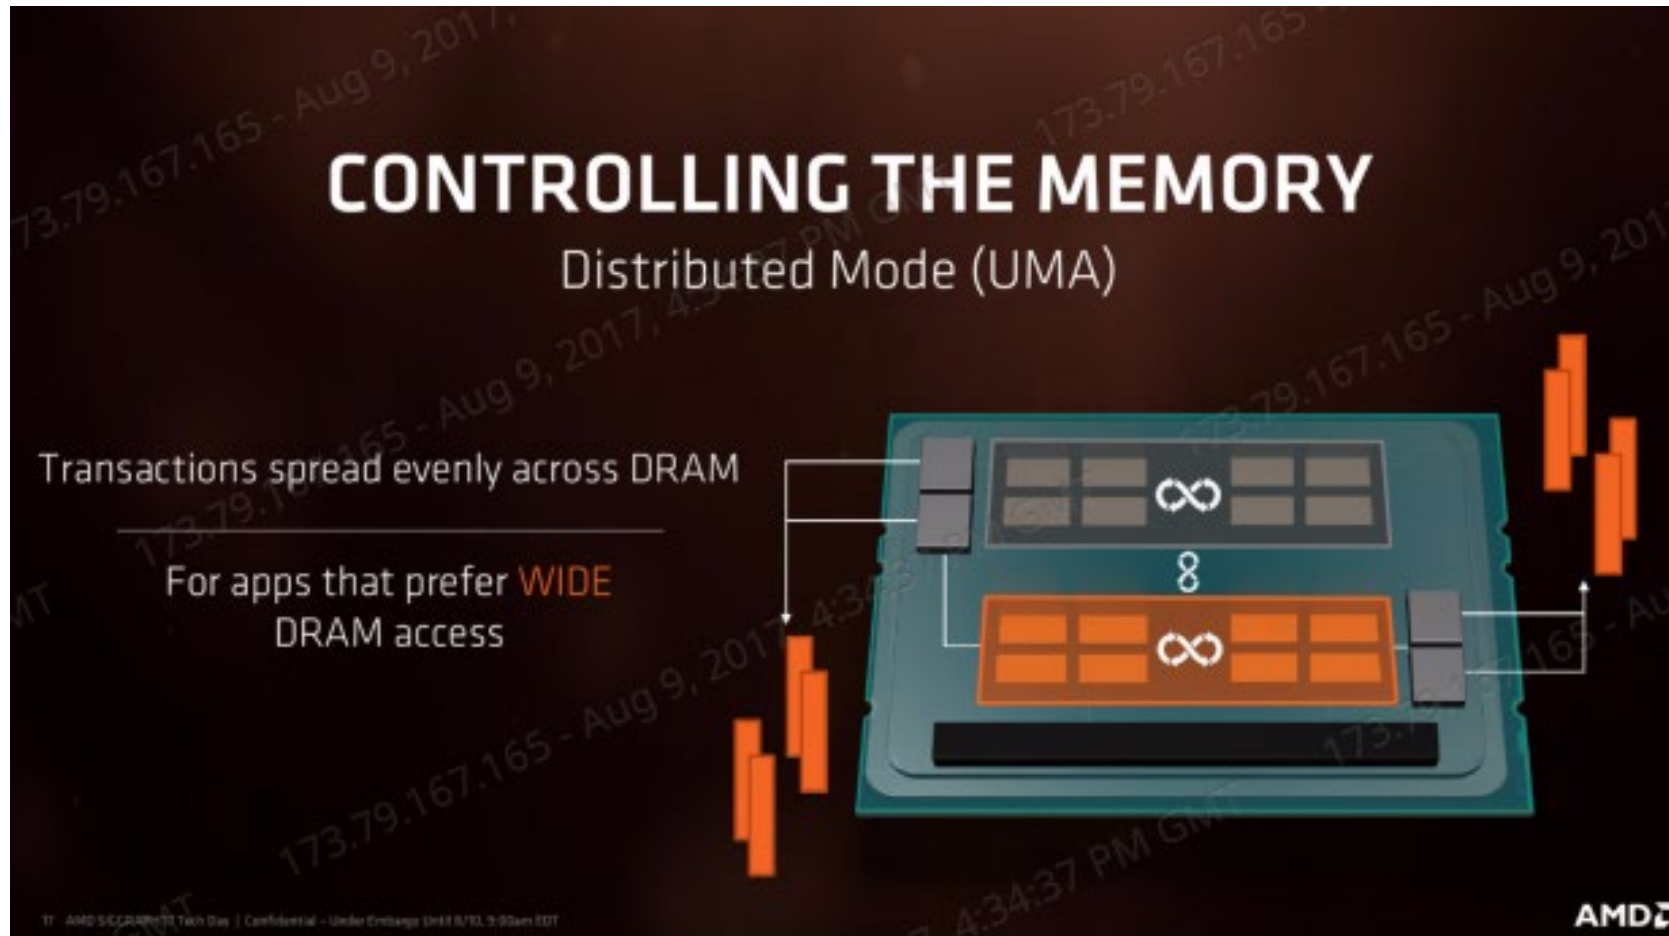
- Latency of accessing the main memory is the same for every processor:
 - Uniform access time, hence the name
- Suitable for small number of processors – due to **contention**
 - Related: Symmetric Multiprocessing (SMP)

Non-Uniform Memory Access (NUMA)

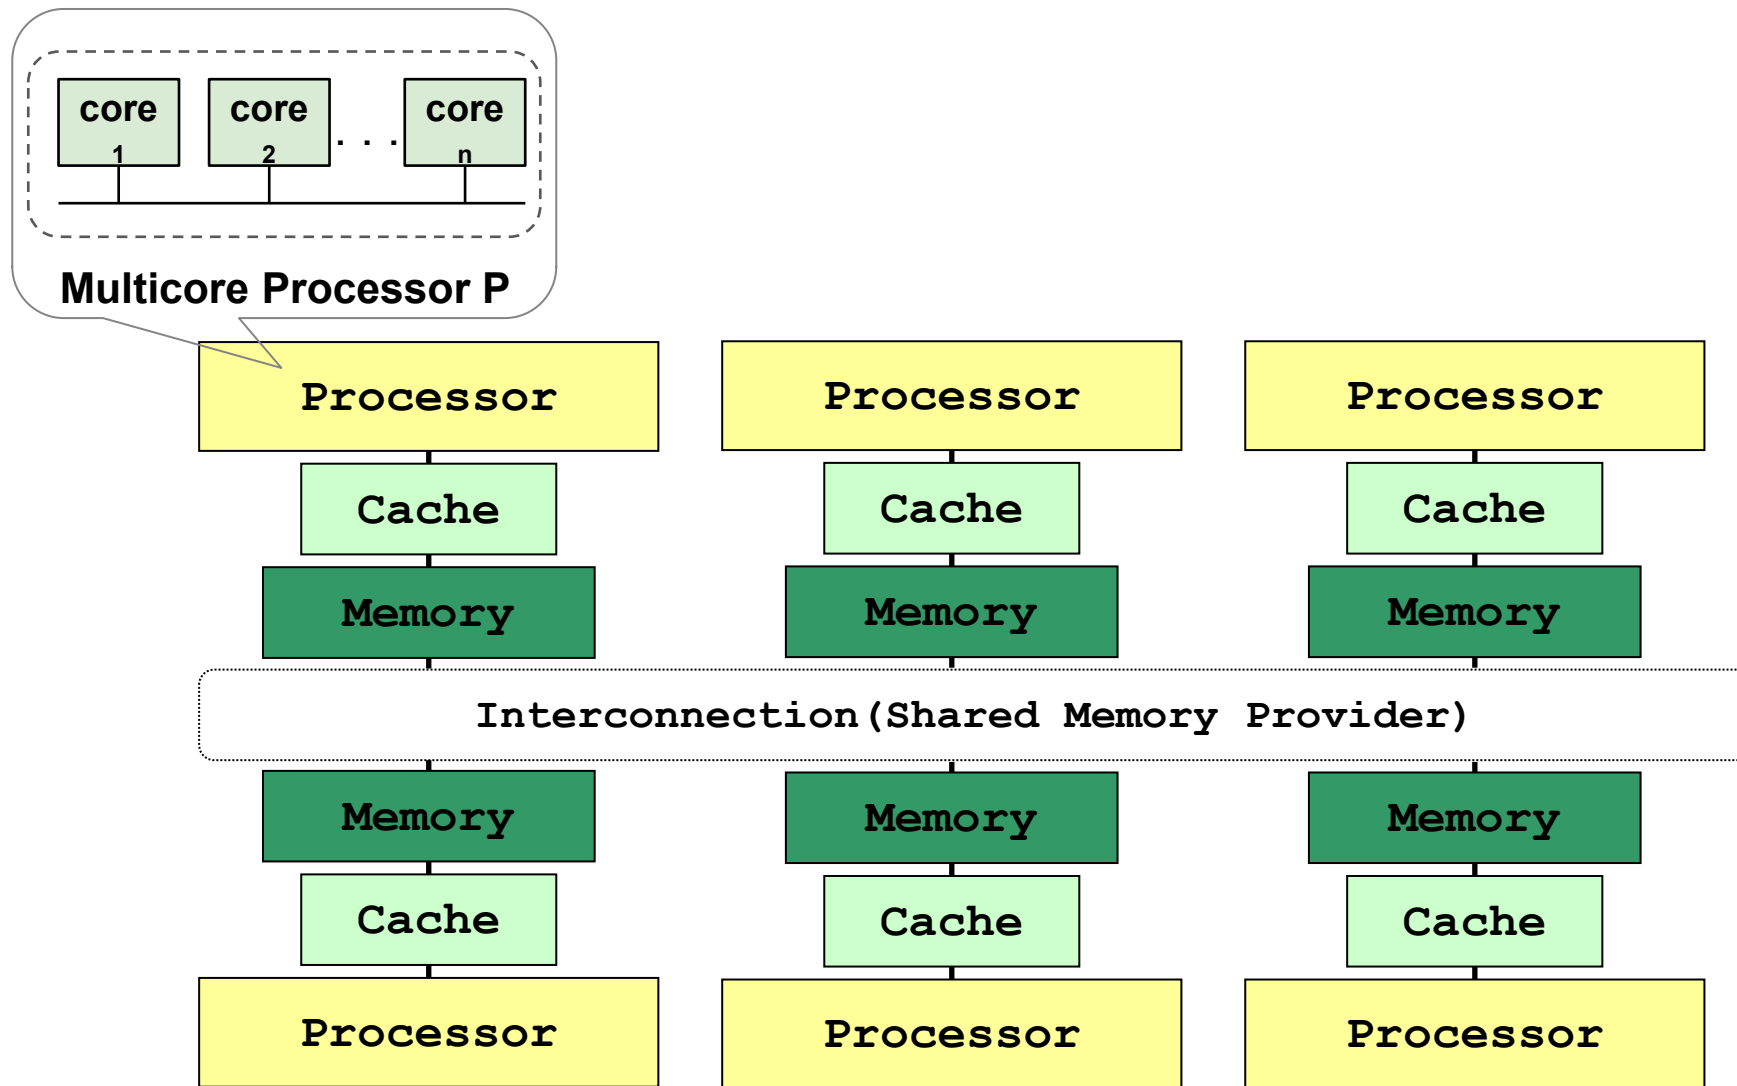


- Physically distributed memory of all processing elements are combined to form a global shared-memory address space:
 - also called **distributed shared-memory**
- Accessing local memory is faster than remote memory for a processor
 - Non-uniform access time


AMD Ryzen



Example: Multicore NUMA

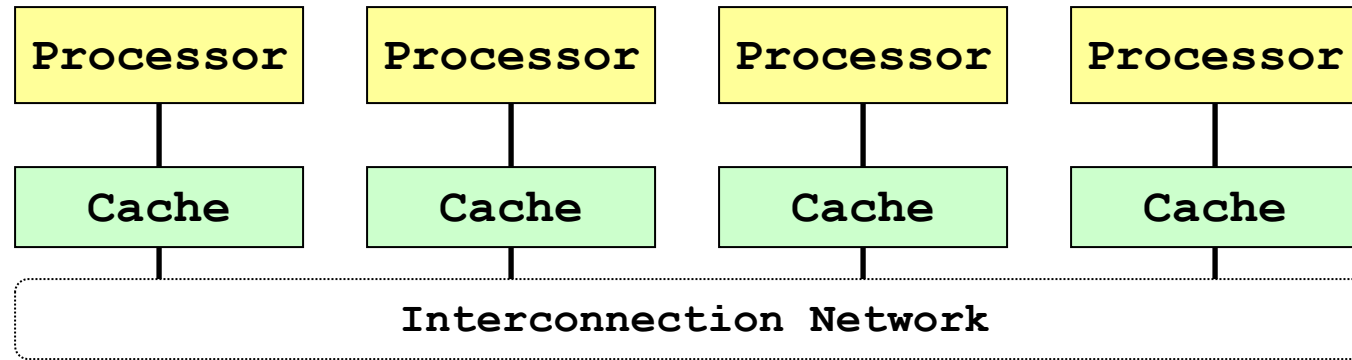


ccNUMA



- Cache Coherent Non-Uniform Memory Access
 - Each node has cache memory to reduce contention

COMA



■ Cache Only Memory Architecture

- ❑ Each memory block works as cache memory
- ❑ Data migrates dynamically and continuously according to the cache coherence scheme

Summary: Shared Memory Systems

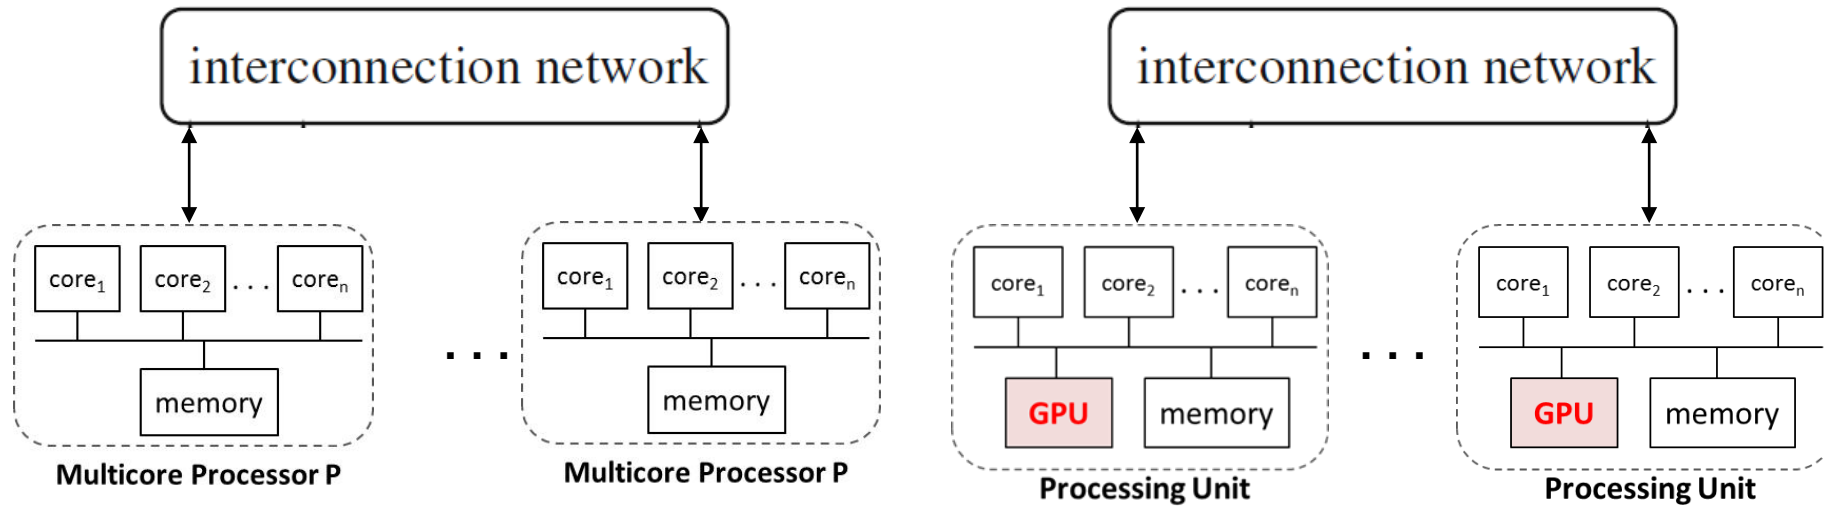

■ Advantages:

- ❑ No need to partition code or data
- ❑ No need to physically move data among processors → communication is efficient

■ Disadvantages:

- ❑ Special synchronization constructs are required
- ❑ Lack of scalability due to **contention**

Hybrid (Distributed-Shared Memory)



**Hybrid with Shared-memory
Multicore Processors**

**Hybrid with Shared-memory Multicore
Processor and Graphics Processing Unit**

Summary

- Goal of parallel architecture is to reduce the average time to execute an instruction
- Various forms of parallelism
- Different types of multicore processors
- Different types of parallel systems and different memory systems