

---

# Parallel Programming Models - II

---

Lecture 09

# Outline

- Data (work) Distribution
  - 1D array
  - 2D array
- Information Exchange
  - Shared variables
  - Communication operations
- Summary


# Data Distribution

- Parallel computing problems are commonly based on array of various dimensions
- Useful to study how to decompose the arrays for distribution on multiple processors
  - known as **data distribution / work distribution / decomposition / partitioning**
- For problems exhibiting data parallelism, data distribution can be used as a simple parallelization strategy

# Data Distribution for 1D Arrays

- Assumptions for discussion:
  - ❑  $p$  identical processors,  $P_1, P_2, \dots, P_p$ , and with processor rank  $i$  in  $\{1, 2, \dots, p\}$
  - ❑ Array elements numbered from 1 to  $n$
- Given a one dimensional array, common distribution patterns:
  - ❑ Blockwise data distribution
  - ❑ Cyclic data distribution


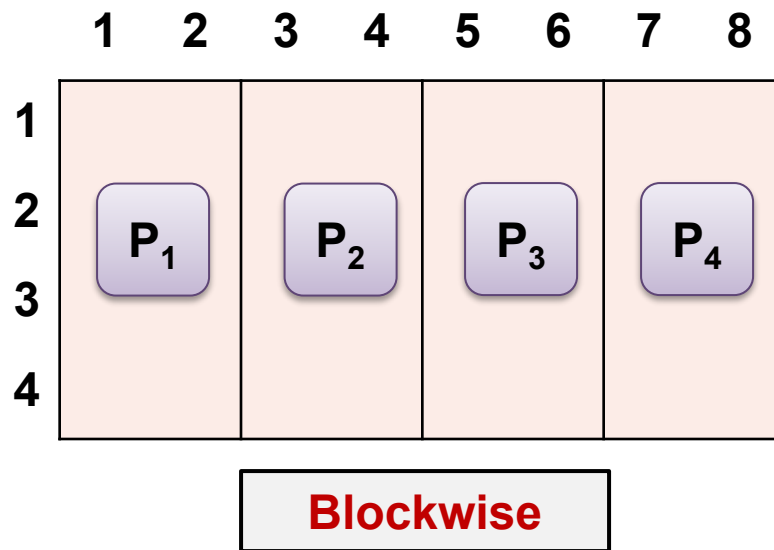
# Blockwise and Cyclic Data Distribution



- Block size,  $B = \left\lceil \frac{n}{p} \right\rceil$
- $P_j$  takes elements  $[(j-1) \times B + 1 \dots j \times B]$
- $P_j$  takes elements  $[j, j+p, \dots, j + (\left\lceil \frac{n}{p} \right\rceil - 1) \times p]$  if  $j \leq n \bmod p$
- $[j, j+p, \dots, j + (\left\lceil \frac{n}{p} \right\rceil - 2) \times p]$  otherwise

# Data distribution for 2D Arrays


- Combination of blockwise / cyclic distribution in one or both dimensions can be used
- One-dimension distributions
  - Use the **column dimension** for illustration:



# Data distribution for 2D Arrays


## ■ One-dimension distributions

- ❑ **Block-Cyclic** is a new distribution pattern
- ❑ Form blocks of size **b**, then perform cyclic (round robin) allocation



# Two Dimensional Distributions

- Processors are **virtually organized into 2D mesh of  $R \times C$** :
  - i.e. each Processor now has a row and column number:



- **Checkerboard** distribution can then be applied:
  - **Blockwise**: elements split into blocks along both dimensions depending on  $R$  and  $C$
  - **Cyclic**: cyclic assignment of elements according to processor mesh
  - **Block-Cyclic**: elements split into  $b_1 \times b_2$  size blocks, then cyclical assignment to processors



# Checkerboard Distribution

	1	2	3	4	5	6	7	8
1	$P_1$				$P_2$			
2								
3	$P_3$				$P_4$			
4								

**Blockwise**

	1	2	3	4	5	6	7	8
1	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$
2	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$
3	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$
4	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$

**Cyclic**

	1	2	3	4	5	6	7	8	9	10	11	12
1	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$
2	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$	$P_1$	$P_2$
3	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$
4	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$	$P_3$	$P_4$

**Block-Cyclic with  $b_1 = 2, b_2 = 2$**

# Exercise: Matrix Multiplication

- To illustrate the effect of data distribution on the computation
- Assume:
  - ❑  $A \times B = C$ , all matrices of  $N \times N$
  - ❑ There are  $p$  processors,  $p$  will be specified
  - ❑ For each value of  $p$ , suggest a data distribution pattern for the matrices  $A$  and  $B$

# Exercise: Matrix Multiplication

1.  $1 < p \leq N$ , you can use  $p = N$  as a start


- ❑ A distributed as
- ❑ B distributed as
- ❑ Each processor calculate

2.  $p = N^2$

- ❑ A distributed as
- ❑ B distributed as
- ❑ Each processor calculate



# Exercise: Heat Transfer Simulation

- A simplistic simulation of heat transfer on a metal plate
- The metal plate is modeled as:
  - 2D integer array
  - Each integer represent the temperature of a "point" on the plate
- The temperature is calculated iteratively:
  - Temperature of a point = Average of the top, left, right and down points



# Exercise: Heat Transfer Simulation

- If we have a  $N \times N$  metal plate and  $p$  processor, where  $p < N$ :
  - Suggest at least two data distribution patterns and discuss pro/cons



I'll trade my B for your A

# INFORMATION EXCHANGE

# Information Exchange

## ■ Purpose

- Information exchange between the executing processors is necessary for controlling the coordination of different parts of a parallel program execution

## ■ Shared address space

- use **Shared variables**

## ■ Distributed address space

- use **Communication operations**

# Shared Variables

- Shared memory programming models assume a global memory accessible by all processors
  - ➔ Information exchange through shared variables
  - ➔ Need synchronization operations for safe concurrent access
- Flow of control abstractions
  - ➔ processes or threads
- Each thread:
  - Executed by one processor or one core in multicore processors
  - Have shared variables and may have private variables



# Synchronized Access

- **Data race:** multiple threads accessing (read and write) the same shared variable
  - Computation result depends on the execution order of threads (**race condition**)
  - May lead to **non-deterministic** behavior
  - Can be avoided using a **critical section** mechanism
- **Critical section:**
  - A program part in which concurrent access should be avoided
    - i.e. only one thread can execute at any point in time
  - Use **mutual exclusion (mutex)** to provide critical section

# Example: OpenMP

- Data race

(might produce a race condition):

```
void main () {  
    int count = 0;  
    #pragma omp parallel  
    {  
        count = count + 1; // race  
    }  
    printf("count = %d\n", count);  
}
```

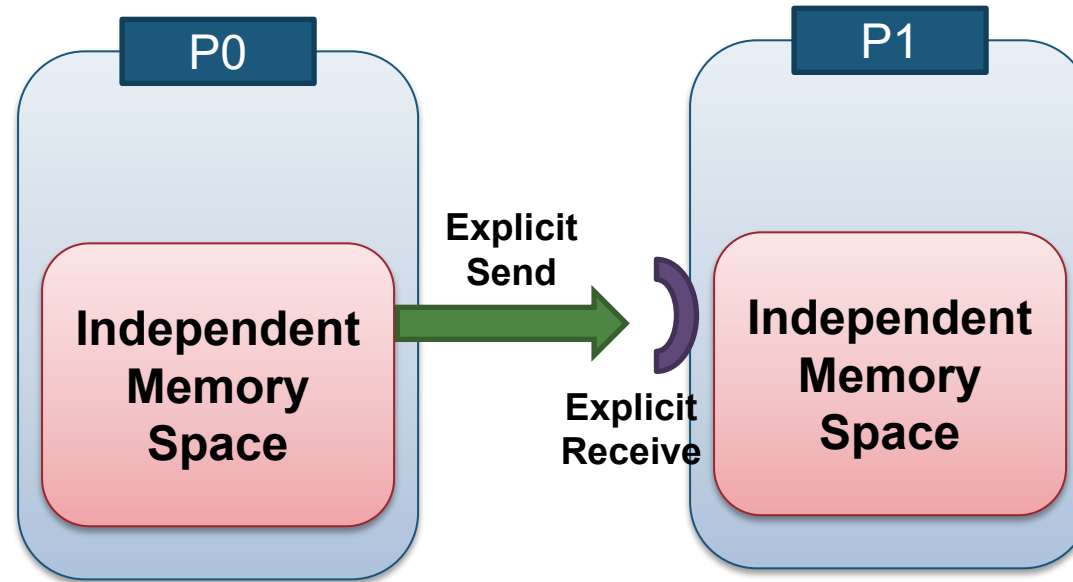
- Mutual exclusion:

```
void main () {  
    int count = 0;  
    omp_lock_t lock;  
    omp_init_lock(&lock);  
  
    #pragma omp parallel  
    {  
        omp_set_lock(&lock);  
        count = count + 1;  
        omp_unset_lock(&lock);  
    }  
    printf("count = %d\n", count);  
}
```

# Communication Operations

- Distributed memory programming models assume disjoint memory space:
  - Exchange of data between processors through **dedicated communication operations**
- One common communication model send / receive messages between participating processors:
  - known as **message-passing programming model**
- Two main types of data exchange:
  - **point-to-point** and global communication

# Principles of Message Passing Model



- Data explicitly partitioned for each process
- All interaction requires both parties to participate
- ➔ The programmer has to explicitly express parallelism

# Principles of Message Passing Model

- **Loosely synchronous** paradigm:
  - Tasks or subsets of tasks synchronize to perform interactions
  - Between these interactions, tasks execute completely **asynchronously**

So, you talk, I talk?

# COMMUNICATION PROTOCOLS

# Send and Receive Operations

```
a = 100;  
send(&a, P1);  
a = 0;
```

Process P0

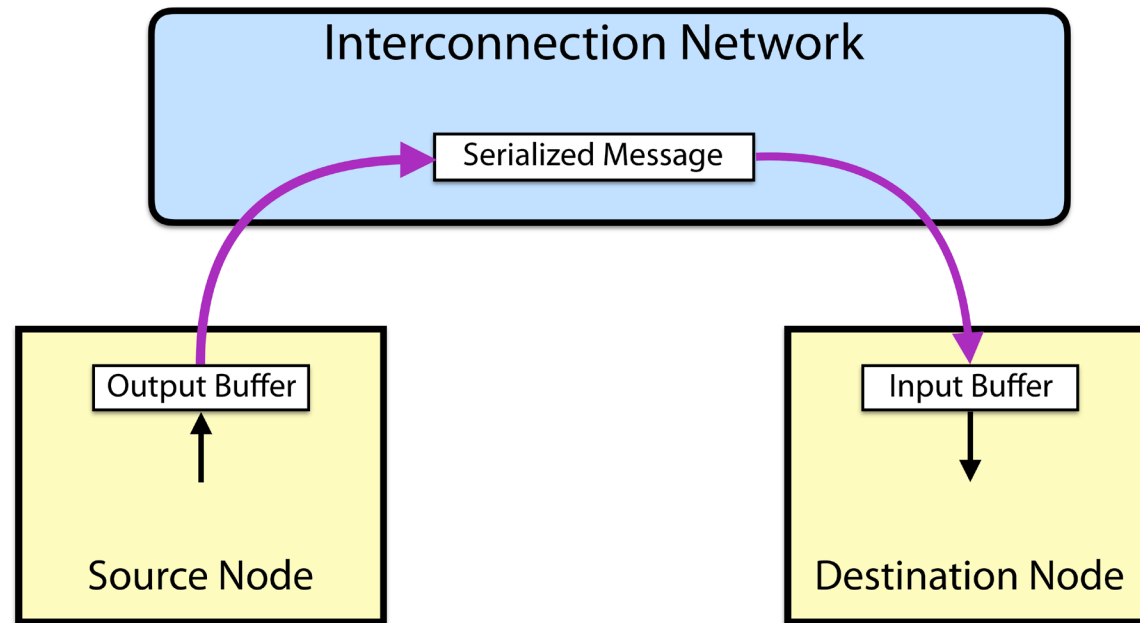
```
receive(&a, P0);  
printf("%d\n", a);
```

Process P1

- Semantic of the send():
  - The value received by P1 should be 100
- Motivates the design of the underlying communication protocols

# Point-to-point Communication (Buffered)

- In a distributed memory system, over a network
  - One-way transfer





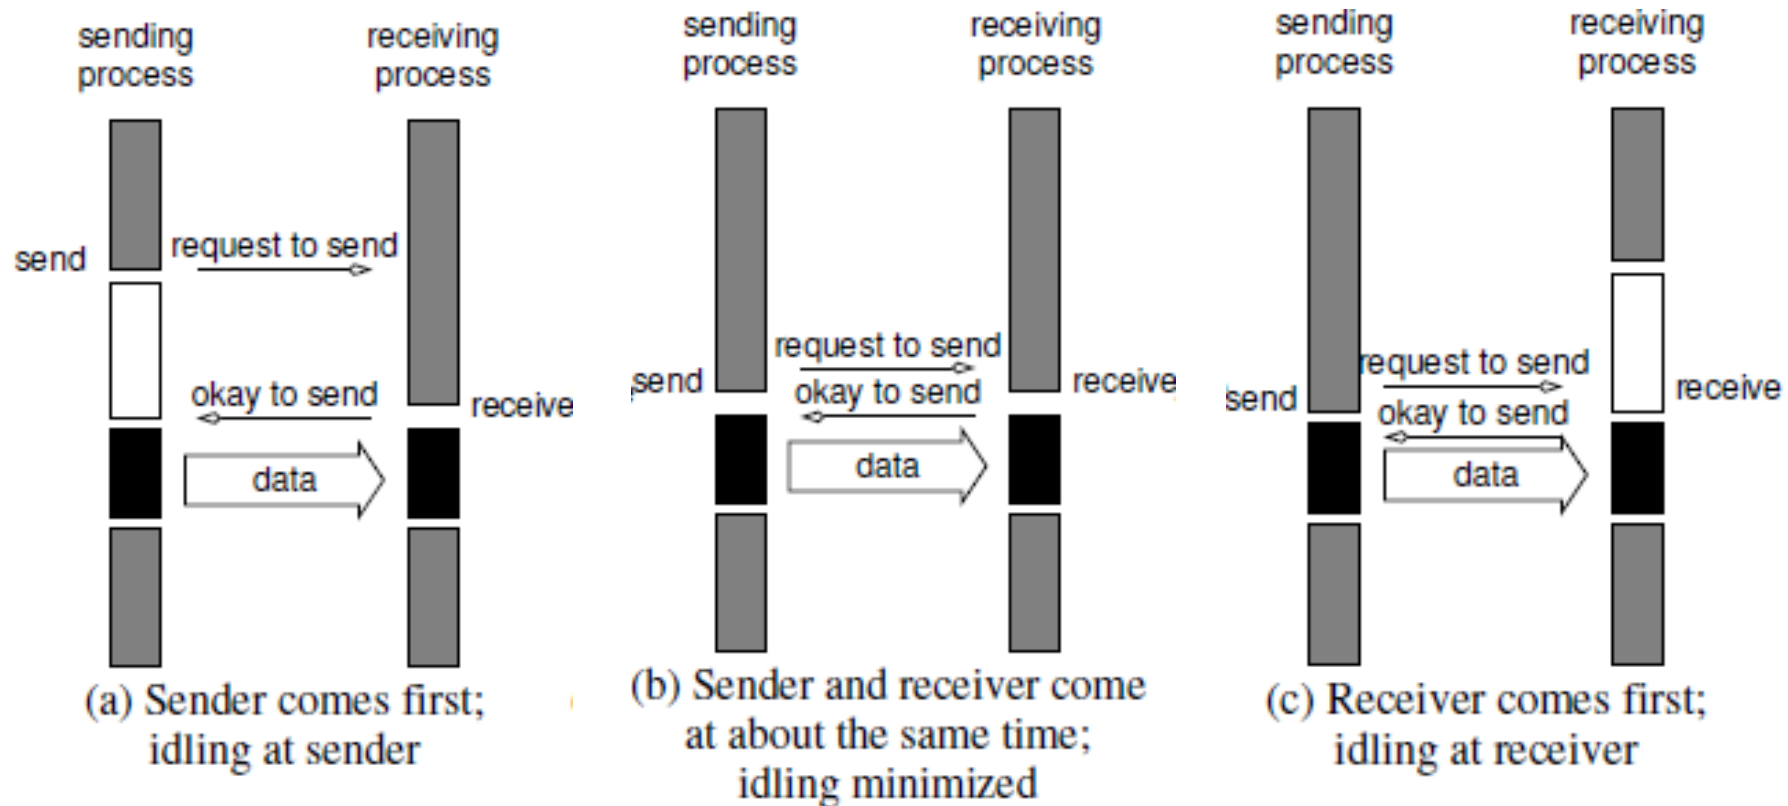
# Send and Receive **Protocols Possibilities**

	<b>Blocking Operations</b>	<b>Non-Blocking Operations</b>
<b>Buffered</b>	Sending process returns after data has been copied into communication buffer	Sending process returns after initiating the transfer to buffer. This operation might not be completed on return.
<b>Non-buffered</b>	Sending process blocks until matching receive operation has been encountered.	
	Send and receives semantics assured by corresponding operation.	Programmer must explicitly ensure completion of the operation by polling.

# Blocking Send

- Send operation blocks until it is **safe** to reuse the input buffer
  - "Safe" refers to the integrity of the data to be sent
- Non-buffered blocking send:
  - The operation blocks until the matching receive has been performed by the receiving process
  - Idling and deadlocks are major issues with non-buffered blocking sends

# Non-Buffered + Blocking Operations

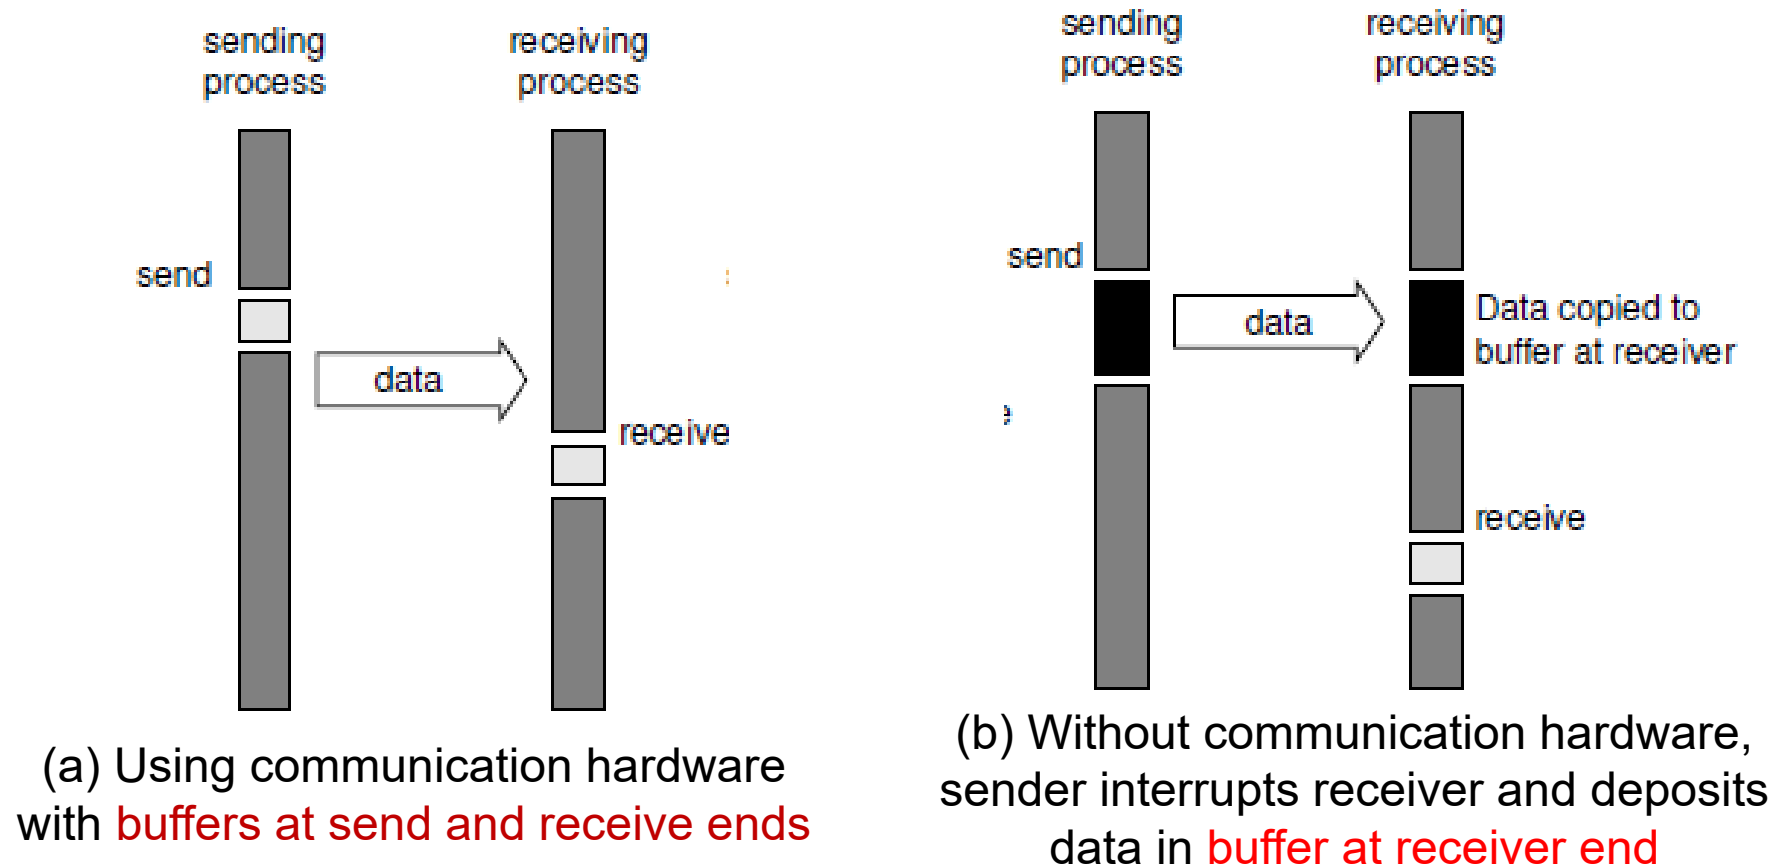


- Considerable idling overheads
  - Due to the mismatch in timing between sender and receiver

# Buffered + Blocking Operations

- To reduce idling overhead:
  - Utilize **buffers** at both ends
- **Sender** simply copies the data into the designated buffer and **returns** after the copy operation has been completed
- **Receiver** similarly buffered the incoming data
- Buffering trades off **idling overhead** for **buffer copying overhead**

# Buffered + Blocking Operations



## Blocking buffered transfer protocols

# Bounded Buffer Size: Impact

```
for (i = 0; i < 1000; i++) {  
    produce(&a);  
    send(&a, P1);  
}
```

Process P0

```
for (i = 0; i < 1000; i++) {  
    receive(&a, P0);  
    consume(&a);  
}
```

Process P1

- What if consumer was much slower than producer?
  - Think "behind the scene"....

# Deadlock

- Deadlocks are still possible with buffering since receive operations block:

```
receive (&a, P1) ;  
send (&b, P1) ;
```

Process P0

```
receive (&a, P0) ;  
send (&b, P0) ;
```

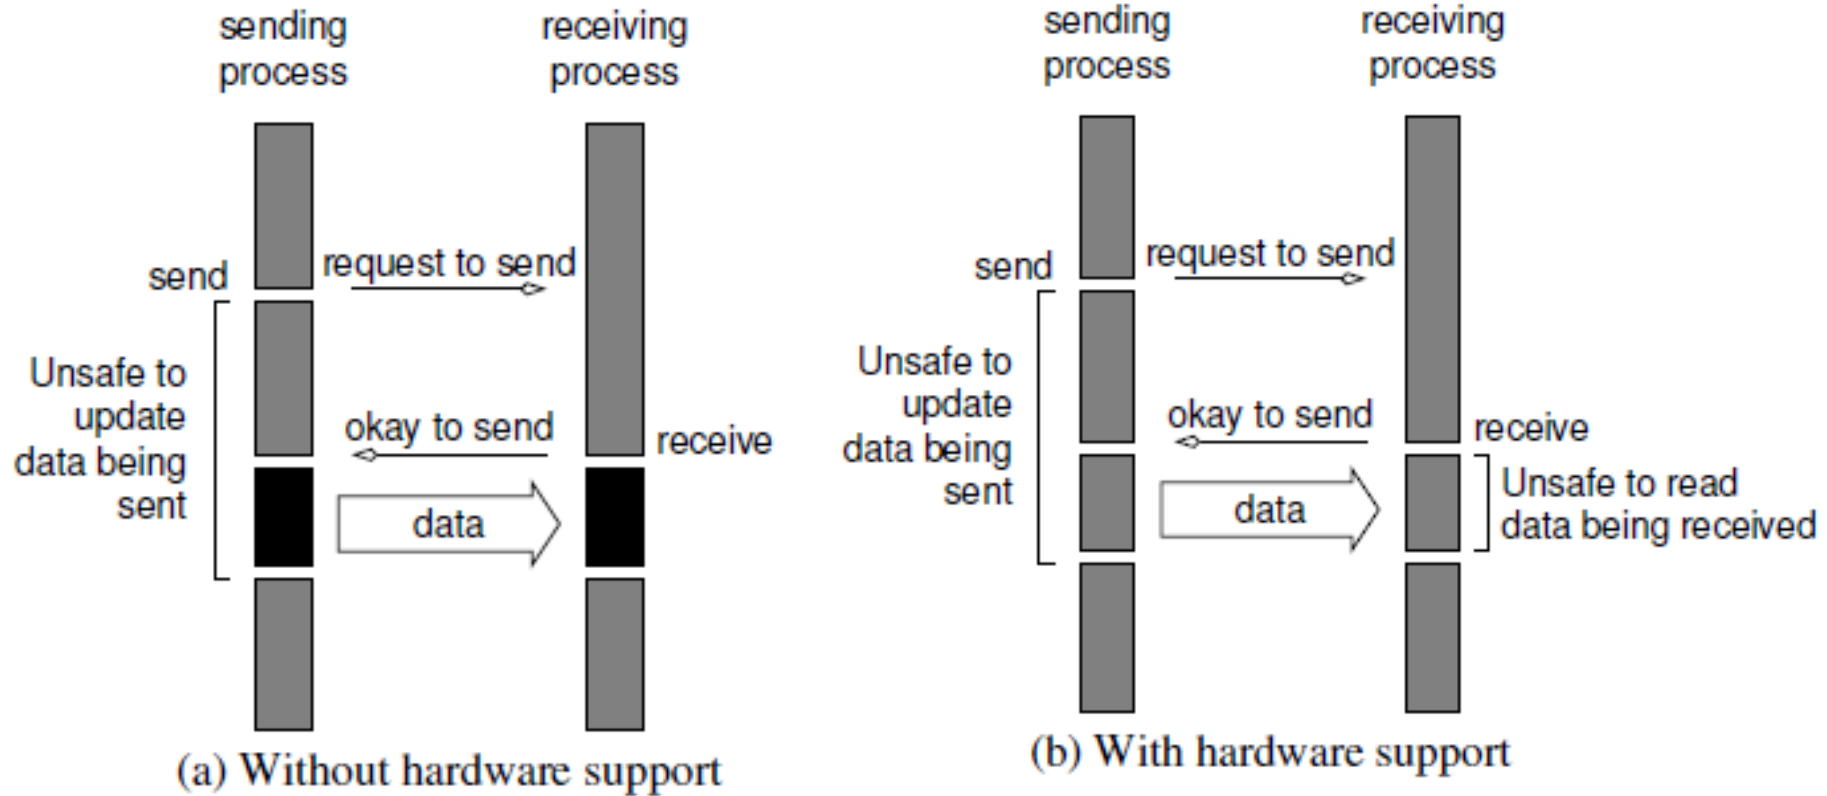
Process P1

# Non-Blocking Operations

- Send / Receive returns before it is semantically safe to use the data transferred
  - Non-blocking operations are generally accompanied by a ***check-status*** operation
  - The programmer must ensure the semantics of the operations
- When used correctly, these primitives are capable of overlapping communication overheads with useful computations
- Message passing libraries typically provide both blocking and non-blocking primitives



# Non-Blocking + Non-Buffered Operations



**Non-blocking non-buffered send and receive operations**

# Semantics of Send/Receive Operations

## Local view

### Blocking

Return from a library call indicates the user is allowed to reuse resources specified in the call

### Non-blocking

A procedure may return before the operation completes, and before the user is allowed to reuse resources specified in the call

## Global view

### Synchronous

Communication operation does not complete before both processes have started their communication operation

### Asynchronous

Sender can execute its communication operation without any coordination with the receiver

# Implementation Options

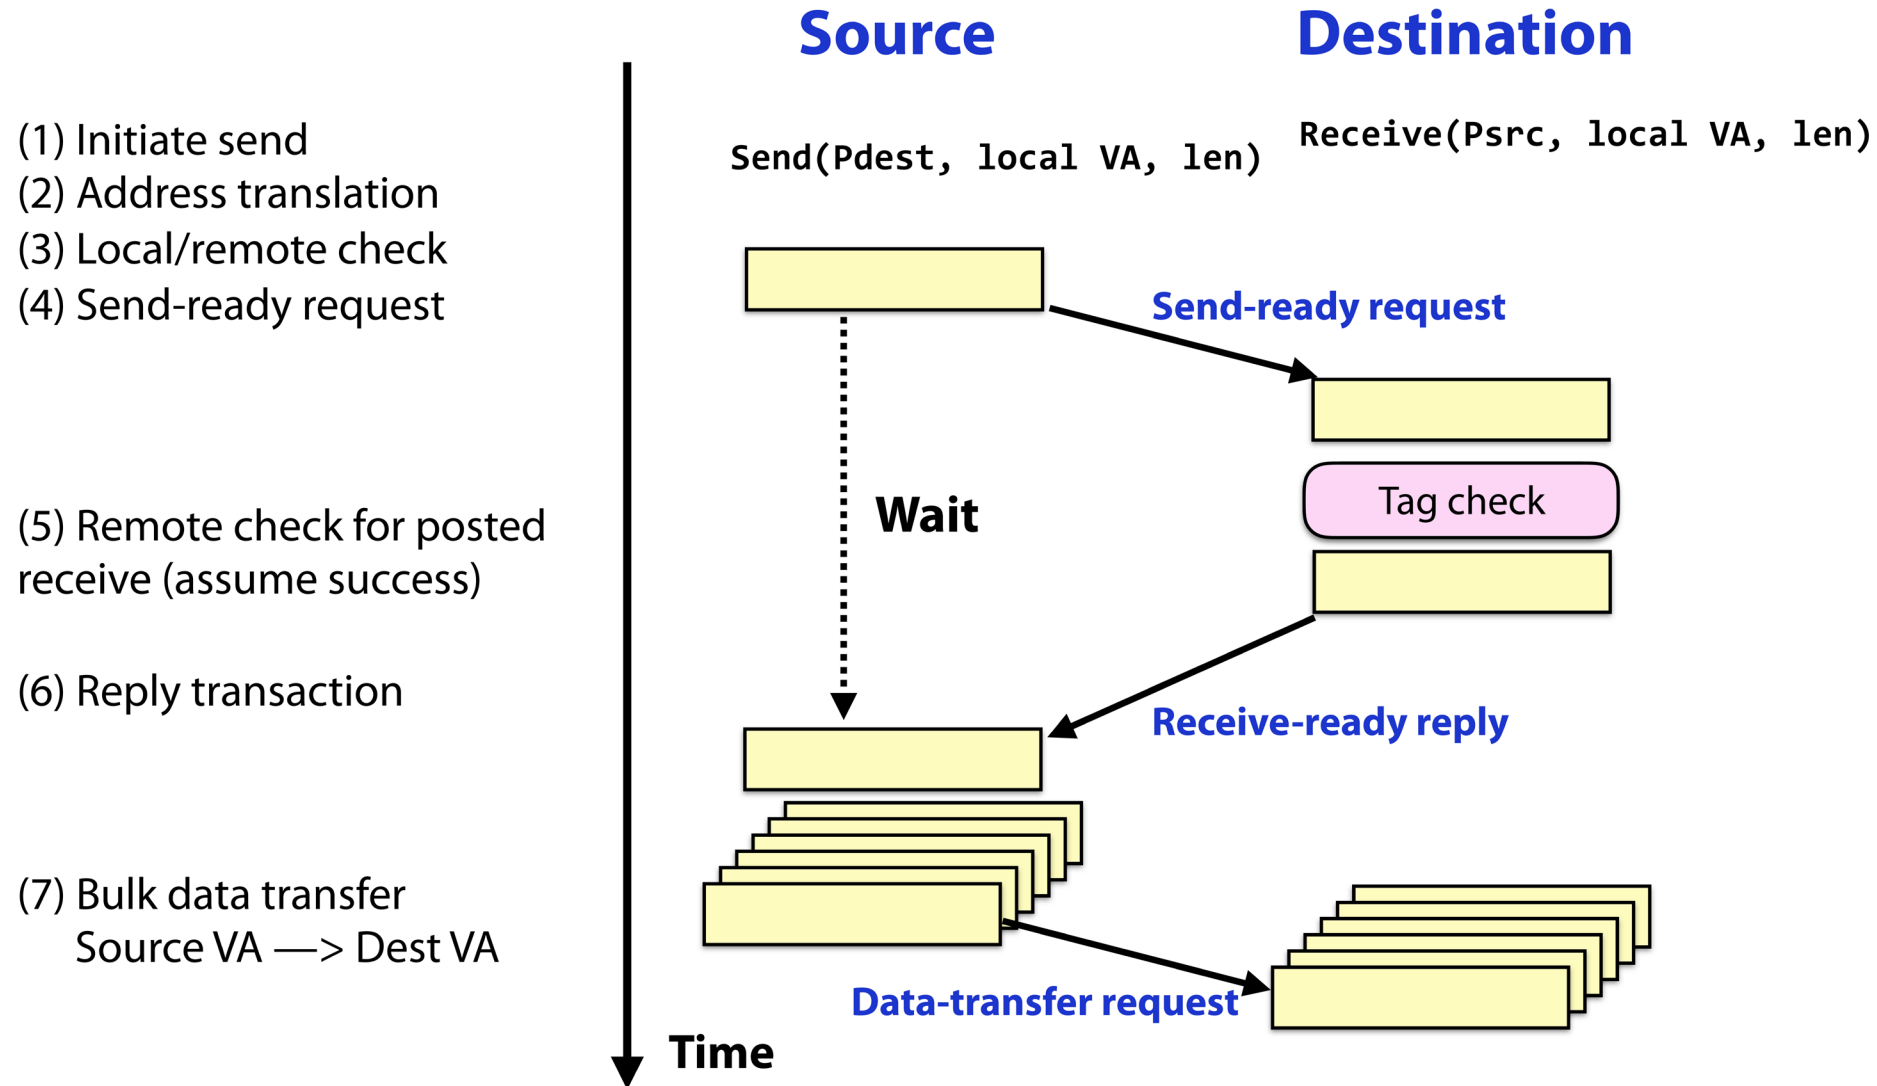
- Synchronous:

- Send completes after matching receive and source data sent
  - Receive completes after data transfer completed from matching send

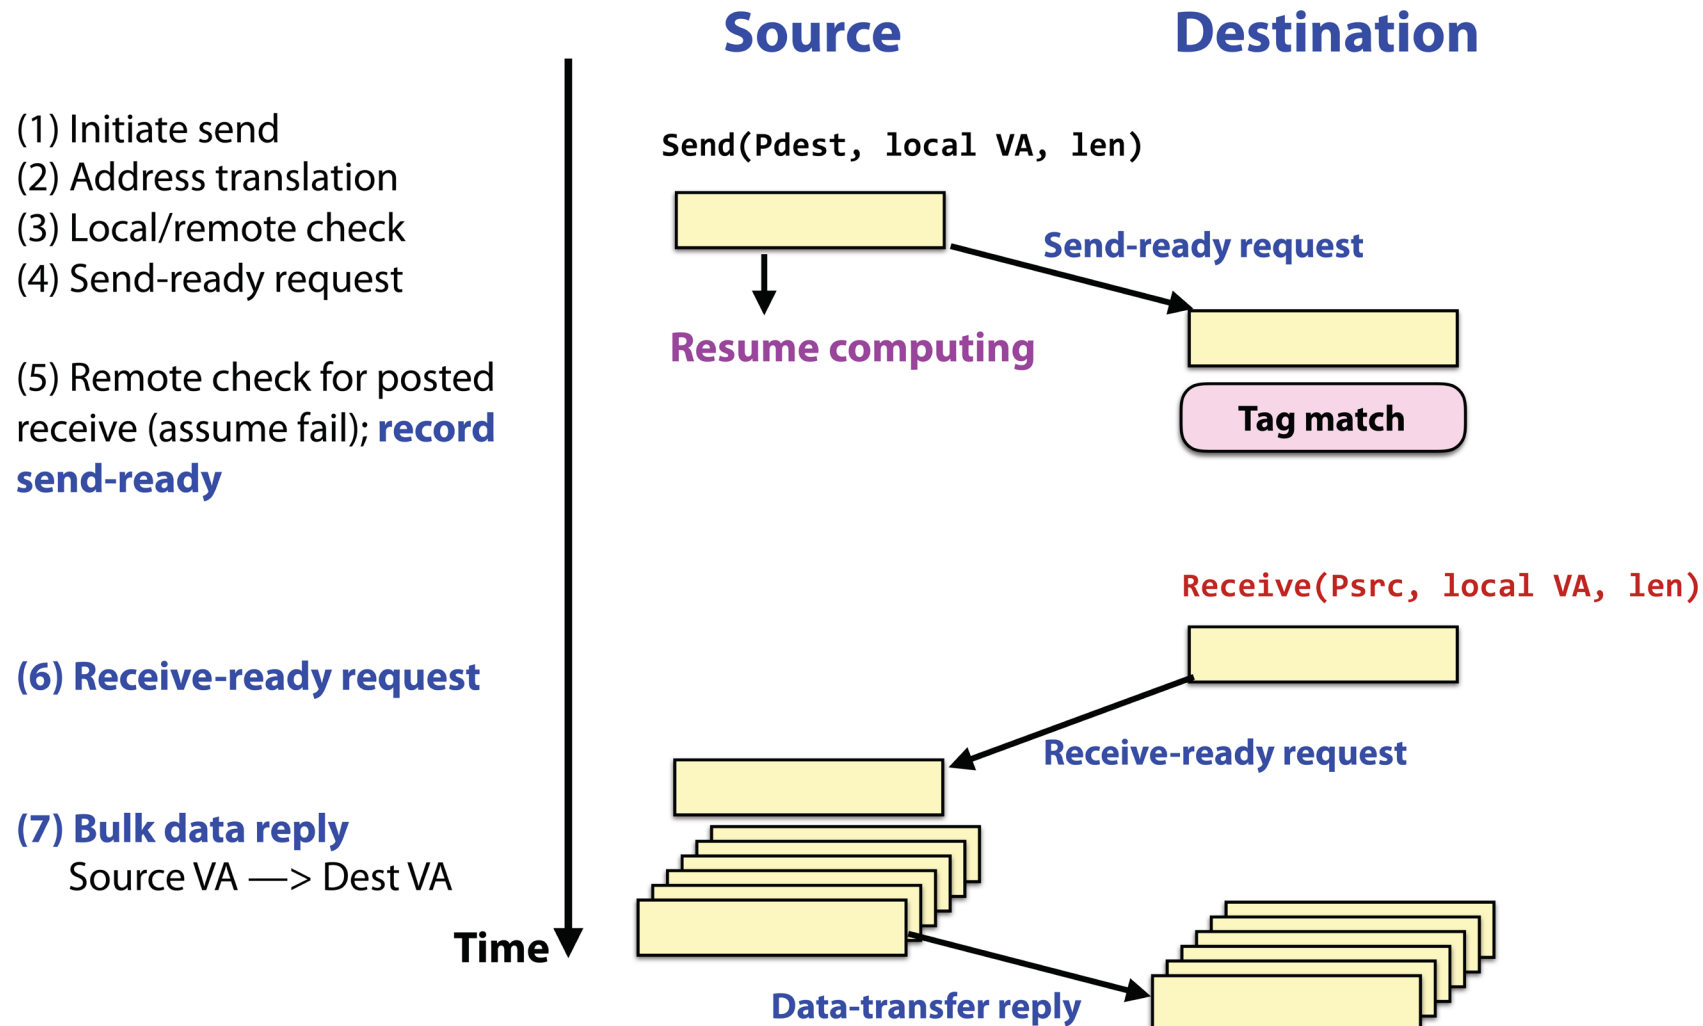
- Asynchronous:

- Send completes after input buffer may be reused

# Synchronous Communication



# Asynchronous Communication



# Summary of Communication Protocols

## 1. **Overhead** of guaranteeing semantic correctness:

- ❑ **Idling** (non-buffered)
- ❑ **Buffer management** (buffered)

## 2. **Side effect:**

- ❑ Safe and easier programming (**blocking**)
- ❑ Hide communication overhead (**non-blocking**)

## 3. **Local or global view:**

- ❑ Synchronous vs. asynchronous communication

# Summary

- Data distribution
- Information exchange for shared and distributed address space