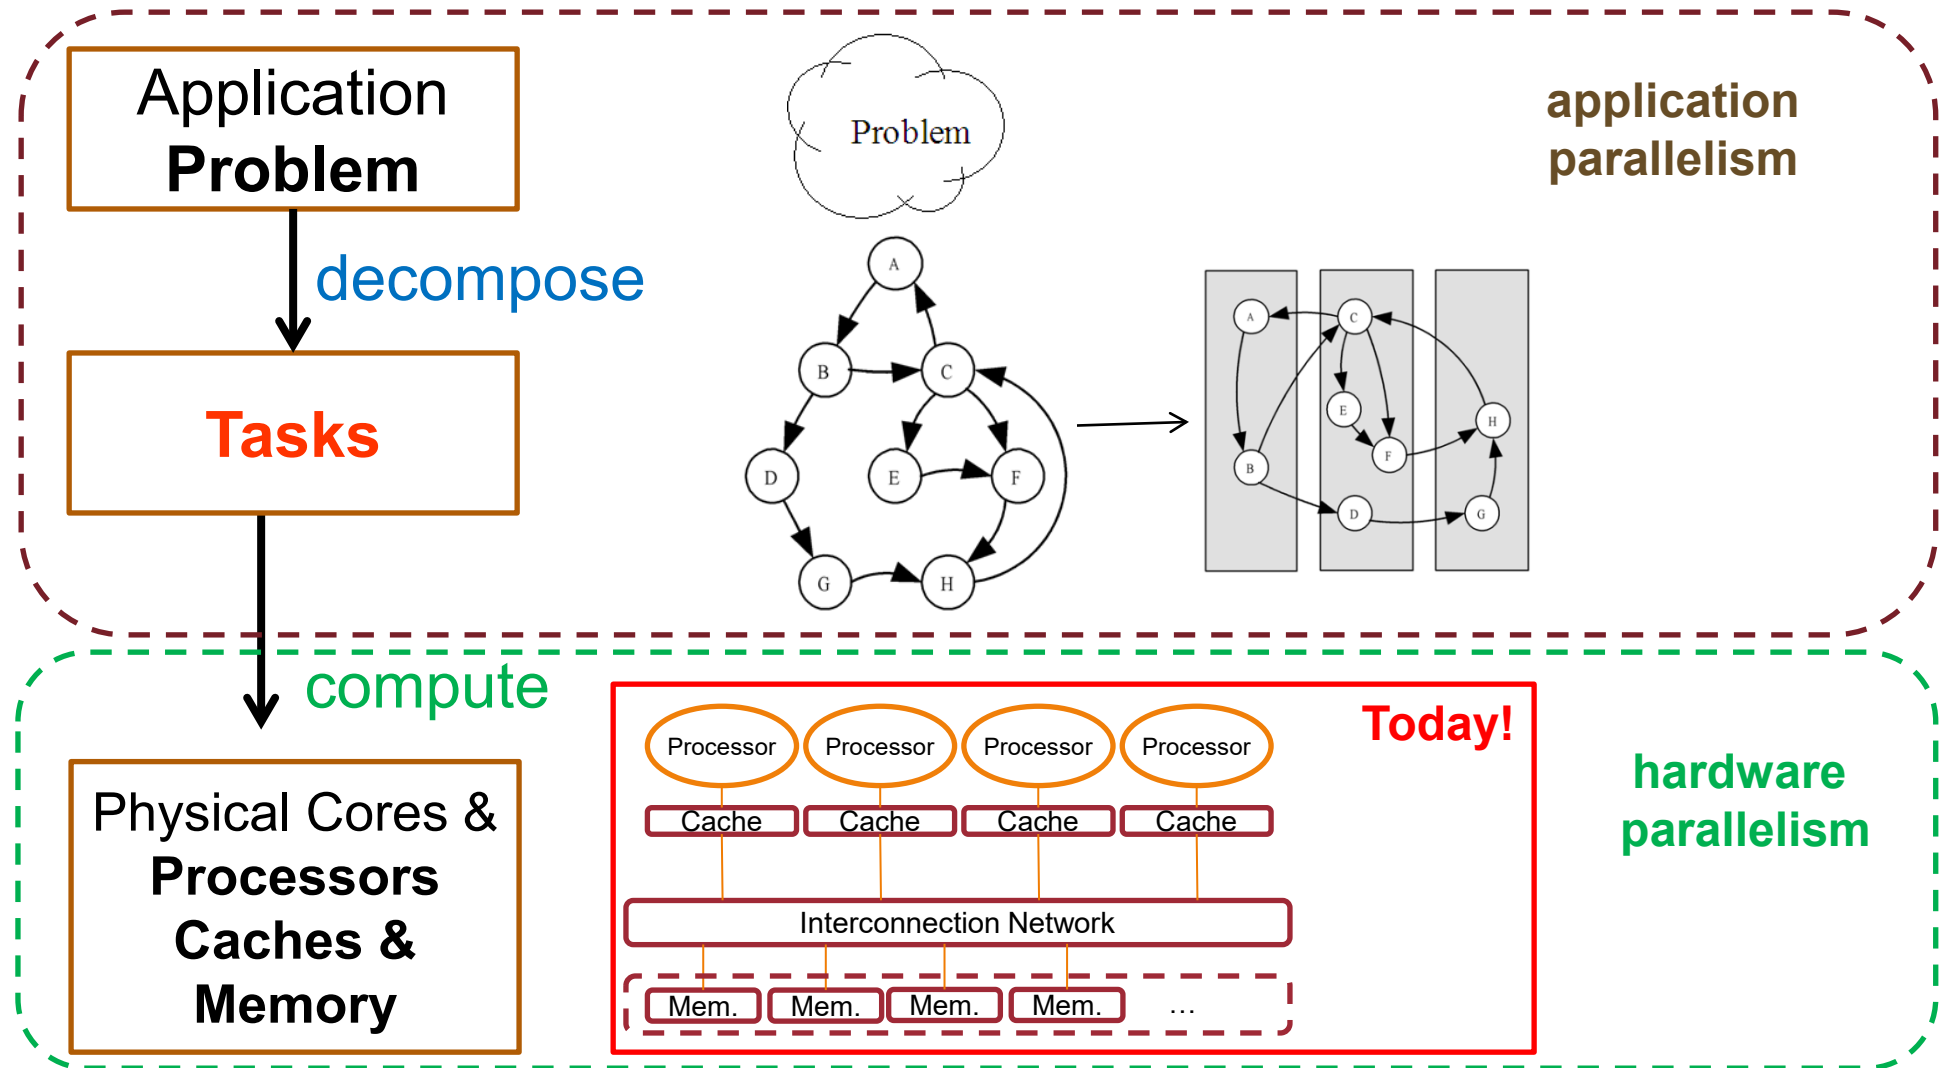


Interconnection Networks

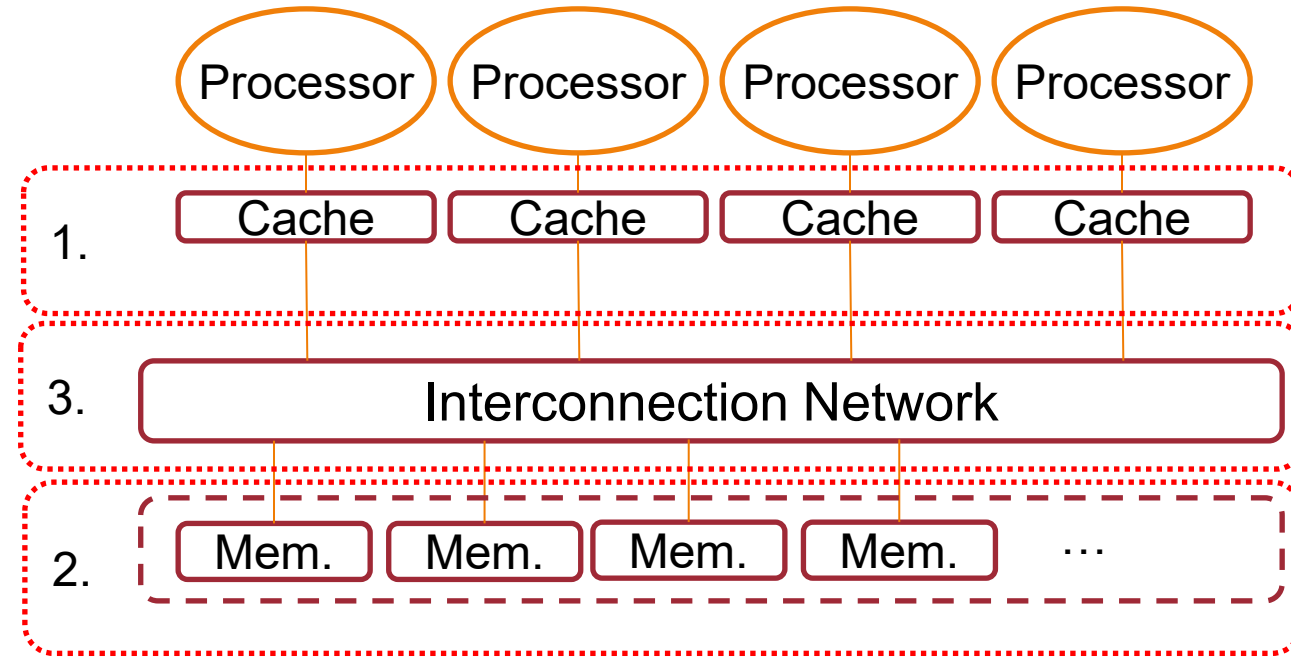
Lecture 11

Parallel Computing




Outline

1. Cache Coherence
2. Memory Consistency
3. Interconnection Networks
 - ❑ Motivating Example
 - ❑ Topology
 - Metrics
 - Direct and Indirect Interconnection Networks
 - ❑ Routing
 - ❑ Current Trends



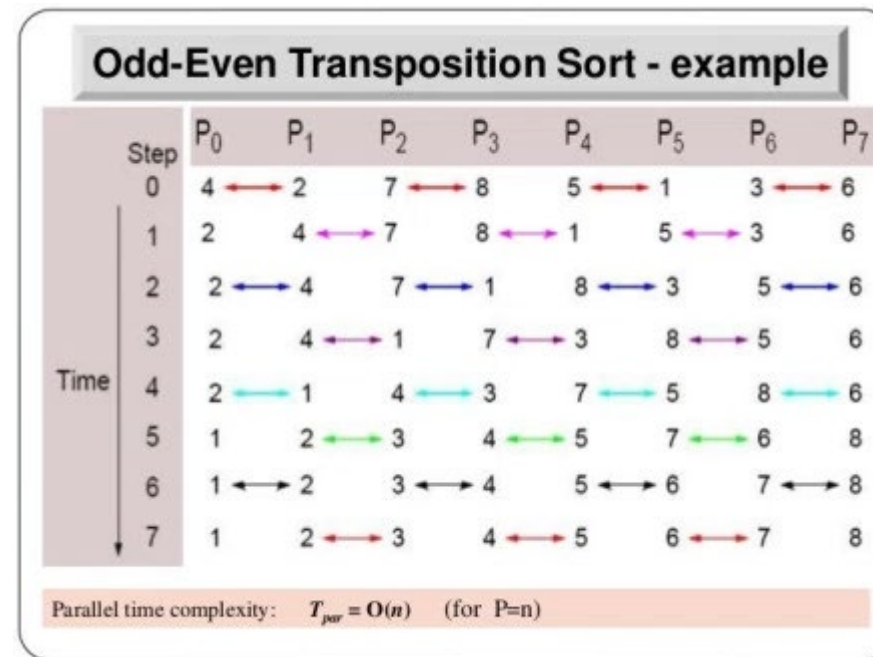
Interconnections: Motivating Example



- Interconnection forms the backbone of communication between:
 - ❑ Processors
 - ❑ Processor and memories
 - ❑ Processors and caches
 - ❑ I/O devices
- Let us first see two examples of interconnect network:
 - ❑ Using simple sorting as the focus point

Sorting on Linear Array


Sort N numbers on N -PEs Linear Array



■ Questions:

- ❑ What is the number rounds needed?
- ❑ Compare with sequential sorting algorithms?

Two Dimensional Mesh




3 x 3 PEs in a **Mesh**

- Observe that the PEs have different number of communication links:
 - PE at the corners = 2 links
 - PE on the edges = 3 links
- If we wrap the left to right; top to bottom
 - ➔ All PEs have four links ➔ known as **Torus**

Algorithm Essential

Sort N elements on N -PEs 2D Mesh

- Each PE has one number initially
- N -PEs arranged in a 2D Mesh:
 - \sqrt{N} rows
 - \sqrt{N} columns
- Sorted into “Snake-Like” order in the end:



Shear Sort Algorithm: Basic Idea

Phase 1: Row Sorting

- Odd Rows sort in **ascending** order
- Even Rows sort in **descending** order


Phase 2: Column Sorting

- All columns sort in **ascending** order (top to bottom)

Repeat


- Until sorted

Example: 16 Numbers – Initial Placement





- Initial placement of 16 numbers (0 to 15)

Phase 1: Row Sort




- Note the different order for each row



Phase 2: Column Sort




Phase 3: Row Sort



Phase 4: Column Sort



Phase 5: Row Sort (Done!)



- Sorted! Note the snake-like arrangement

Complexity

- A bold statement:

For N Numbers, **$\text{Log}_2 N + 1$** phases

- E.g. 16 Numbers = $\text{Log}_2 16 + 1$
- = 5 phases
- How to perform each of the row/column sort?
 - Use ... odd-even transposition sort
 - So, complexity is



We now know that

- Connection "patterns" enable different algorithms
- Connection "patterns" have a huge impact on execution
- Follow up questions:
 - What are the common ways to connect?
 - How to handle communication in these interconnection networks?

Interconnection Network Impact

- System scalability – size and extensibility
- System performance and energy efficiency
 - Communication speed
 - Latency to memory
 - Energy spent to communicate

Major Questions in Interconnections

Topology

- What is the geometrical shape of the connection?

Routing

- What is the path for a message to follow?

Switching

- How to transfer a message along a path?


Flow Control

- How to handle concurrent messaging?

Beautiful patterns that induce headache....


TOPOLOGY

Topology: Major Type



Direct Interconnection

- Also known as **Static** or **Point-to-Point**
- Usually endpoints are of the same type (core, memory)



Indirect Interconnection

- Also known as Dynamic
- Interconnect is formed by switches

Direct Interconnection: Overview

- Topology can be modeled as a Graph:
 - $G = (V, E)$, V = vertices; E = edges
 - Can use graph theory to understand!
- However, some of the general cases are very hard to determine
 - Only study specific subcases
- Metrics:
 - Diameter
 - Degree
 - Bisection Width
 - Connectivity

Topology Metric - Diameter

- **Diameter $\delta(G)$:** **maximum distance** between any pair of nodes

$$\delta(G) = \max_{u,v \in V} \min_{\substack{\varphi \text{ path} \\ \text{from } u \text{ to } v}} \{k \mid k \text{ is the length of the path } \varphi \text{ from } u \text{ to } v\}.$$

- **Usefulness:**
 - Small diameter ensures small distances for message transmission

Topology Metric - Degree

- **Degree $g(v)$:** number of direct neighbour nodes of node v
 - Degree $g(G)$: maximum degree of a node in a network G

$$g(G) = \max\{g(v) \mid g(v) \text{ degree of } v \in V\}.$$

- **Usefulness:**
 - Small node degree reduces the node hardware overhead

Topology Metric - Bisection Width

- **Bisection width $B(G)$:** minimum number of edges that must be removed in to divide the network into two equal halves

$$B(G) = \min_{\substack{U_1, U_2 \text{ partition of } V \\ ||U_1| - |U_2|| \leq 1}} |\{(u, v) \in E \mid u \in U_1, v \in U_2\}|.$$

- Bisection bandwidth $BW(G)$: Total bandwidth available between the two bisected portion of the network
- **Usefulness:** a measure for the capacity of a network when transmitting messages simultaneously

Topology Metric - Connectivity

- **Node connectivity $nc(G)$:** minimum number of nodes that must fail to disconnect the network

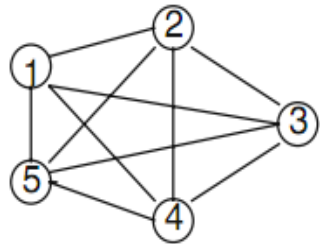
$$nc(G) = \min_{M \subset V} \{ |M| \mid \text{there exist } u, v \in V \setminus M, \text{ such that there exists no path in } G_{V \setminus M} \text{ from } u \text{ to } v \}.$$

- **Usefulness:** Determine the robustness of the network
- **Edge connectivity $ec(G)$:** minimum number of edges that must fail to disconnect the network

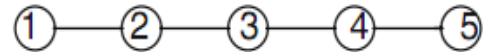
$$ec(G) = \min_{F \subset E} \{ |F| \mid \text{there exist } u, v \in V, \text{ such that there exists no path in } G_{E \setminus F} \text{ from } u \text{ to } v \}.$$

- **Usefulness:** Determine number of independent paths between any pair of nodes

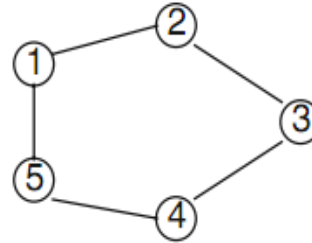
Examples



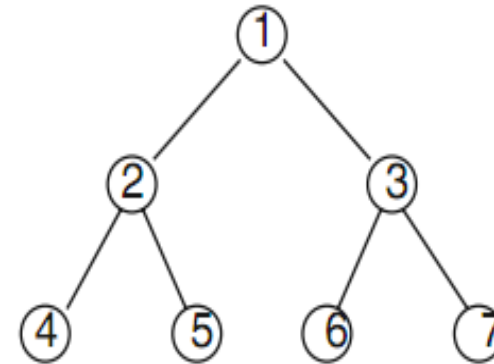
Complete network



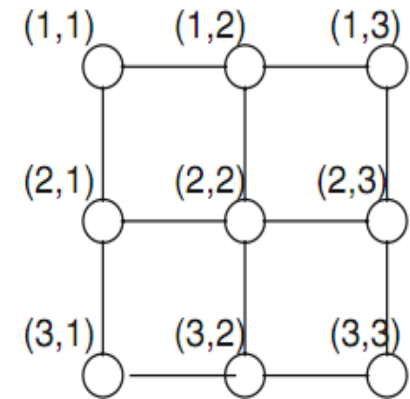
Linear array network




Ring network





complete binary tree



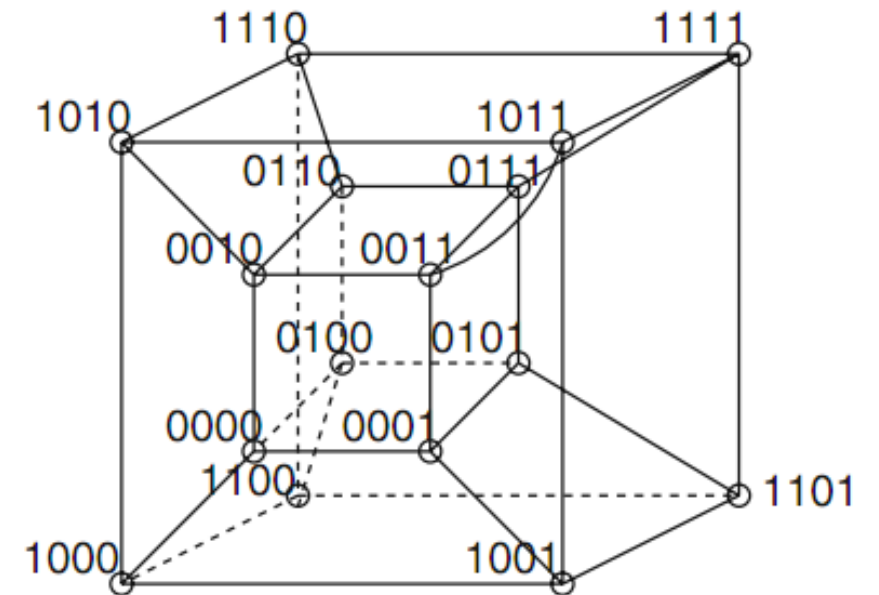
2-dimensional mesh




2-dimensional torus




Hypercube




Complete, Linear Array & Ring



Complete network



Linear array network




Ring network


Degree $g(G)$	$n - 1$	2	2
Diameter $\delta(G)$	1	$n - 1$	$\lceil \frac{n}{2} \rceil$
Edge Connectivity $ec(G)$	$n - 1$	1	2
Bisection Width $B(G)$	$\left(\frac{n}{2}\right)^2$	1	2

2-Dimensional Mesh and Torus

$$n = r^2$$



2-dimensional mesh




2-dimensional torus

Degree $g(G)$	4	4
Diameter $\delta(G)$	4	2
Edge Connectivity $ec(G)$	2	4
Bisection Width $B(G)$	4	8 6

■ Generalize to d -dimension

□ Number of node $n = r^d$




Complete Binary Tree




complete binary tree

Degree $g(G)$	3
Diameter $\delta(G)$	4
Edge Connectivity $ec(G)$	1
Bisection Width $B(G)$	1

Hypercube




$$n = 2^k$$




k-dimensional hypercube for k=1,2,3,4

Degree $g(G)$	$\log n$
Diameter $\delta(G)$	$\log n$
Edge Connectivity $ec(G)$	$\log n$
Bisection Width $B(G)$	$n/2$

Cube-Connected-Cycles (CCC)



Hypercube, $k = 3$

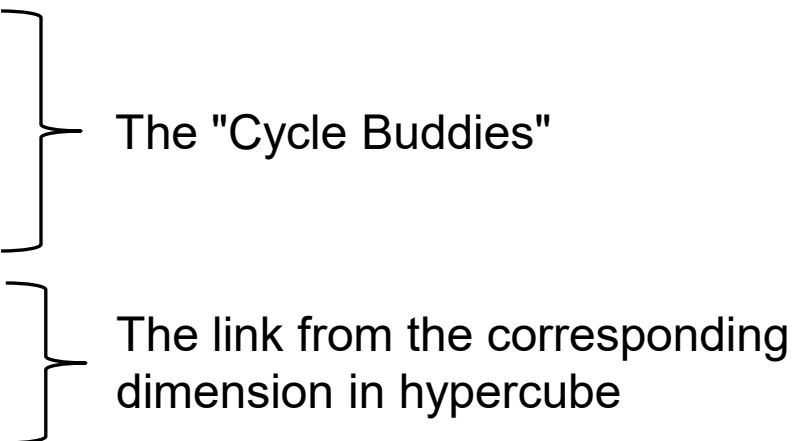


Cube-Connected-Cycles network for $k = 3$

- From a k -dimensional hypercube ($k \geq 3$)
 - Substitute each node with a cycle of k -nodes
 - Each of the k -nodes take one of the original k links
 - ➔ Total nodes = $k2^k$

CCC - Construction


- Each node in a k -dimensional CCC is labeled as (X, Y)
 - X = the corresponding node index in hypercube
 - Y = the position in the cycle, i.e. $0 \dots k-1$
- Node (X, Y) is connected to:
 - $(X, (Y+1) \bmod K)$
 - $(X, (Y-1) \bmod K)$
 - $(X \oplus 2^Y, Y)$



The "Cycle Buddies"

The link from the corresponding dimension in hypercube

CCC - Metrics



Degree $g(G)$	3
Diameter $\delta(G)$	$2k - 1 + \left\lfloor \frac{k}{2} \right\rfloor$
Edge Connectivity $ec(G)$	3
Bisection Width $B(G)$	$\frac{n}{2k}$

Many others

- Generalized d-dimensional mesh
- Generalized d-dimensional torus
- Generalized k-ary d-dimensional cube
- Kautz Graph
- etc....

Summary of Metrics

network G with n nodes	degree $g(G)$	diameter $\delta(G)$	edge- connectivity $ec(G)$	bisection bandwidth $B(G)$
complete graph	$n - 1$	1	$n - 1$	$\left(\frac{n}{2}\right)^2$
linear array	2	$n - 1$	1	1
ring	2	$\left\lfloor \frac{n}{2} \right\rfloor$	2	2
d -dimensional mesh ($n = r^d$)	$2d$	$d(\sqrt[d]{n} - 1)$	d	$n^{\frac{d-1}{d}}$
d -dimensional torus ($n = r^d$)	$2d$	$d \left\lfloor \frac{\sqrt[d]{n}}{2} \right\rfloor$	$2d$	$2n^{\frac{d-1}{d}}$
k -dimensional hyper- cube ($n = 2^k$)	$\log n$	$\log n$	$\log n$	$\frac{n}{2}$
k -dimensional CCC-network ($n = k2^k$ for $k \geq 3$)	3	$2k - 1 + \lfloor k/2 \rfloor$	3	$\frac{n}{2k}$
complete binary tree ($n = 2^k - 1$)	3	$2 \log \frac{n+1}{2}$	1	1
k -ary d -cube ($n = k^d$)	$2d$	$d \left\lfloor \frac{k}{2} \right\rfloor$	$2d$	$2k^{d-1}$

Indirect Interconnection: Overview

■ Why?

- ❑ Reduce hardware costs by sharing switches and links


■ How?

- ❑ Switches provide indirect connection between nodes and can be configured dynamically

■ What metric?

- ❑ Cost (number of switches / links)
- ❑ Concurrent connections

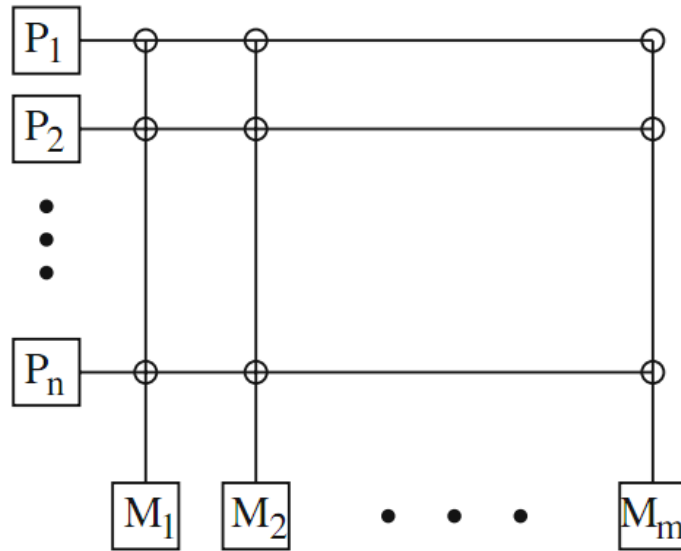
Bus Network




- A set of wires to transport data from a sender to a receiver
- Only one pair of devices can communicate at a time
 - ❑ A bus arbiter is used for the coordination
 - ❑ ➔ Typically used for a small number of processors

Crossbar Network


- A $n \times m$ crossbar network has n inputs and m outputs



$n \times m$ crossbar network



Straight state

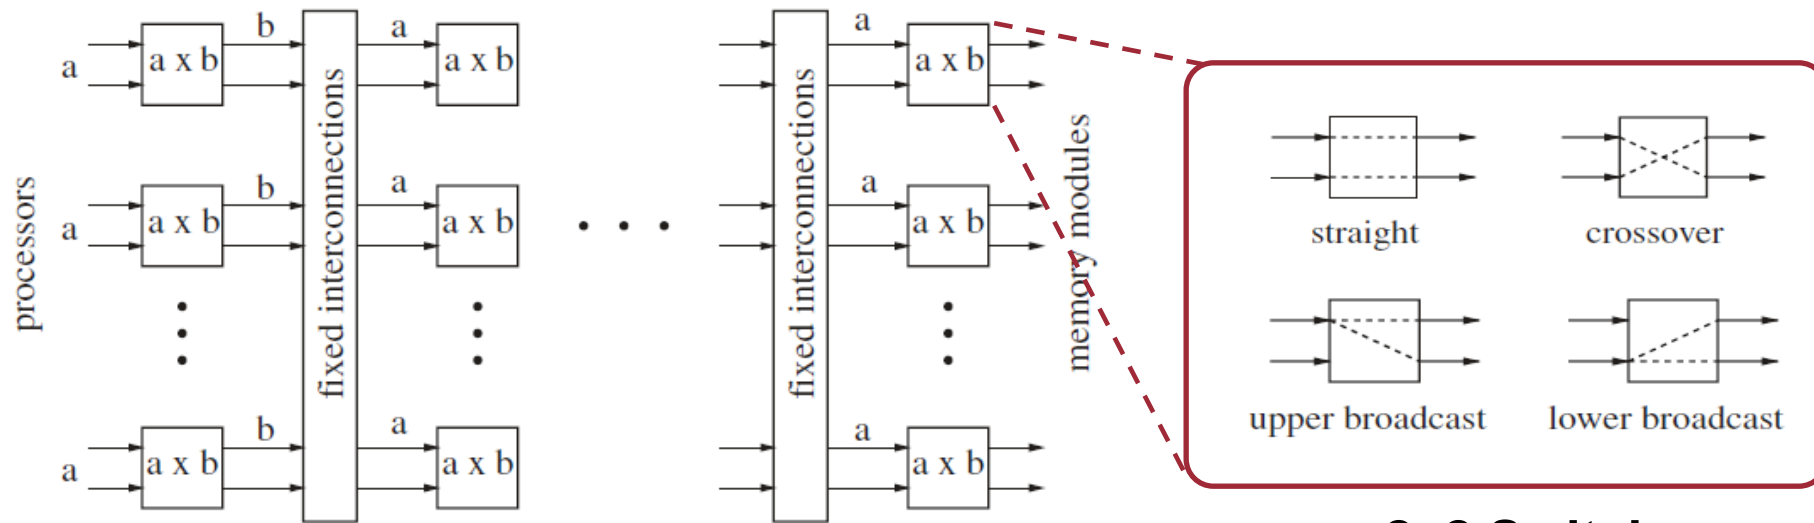


Direction change state

- Two states of a switch: **straight** or **direction change**
- Hardware is costly ($n \times m$ switches) \rightarrow small number of processors

Multistage Switching Network

- Several intermediate switches with connecting wires between neighbouring stages
- **Goal:** obtain a small distance for arbitrary pairs of input and output devices




**Multistage Networks using
Regular 2x2 Crossbar Switches**

**2x2 Switch
Settings**

Omega Network

- One unique path for every input to output




16 × 16 Omega Network

- An $n \times n$ Omega network has **$\log n$ stages**
 - **$n/2$ switches per stage**
 - Connections between stages are regular
 - Also known as $(\lg n - 1)$ – dimension Omega Network

Omega Network - Construction

- A switch position: (α, i)
 - α : position of a switch within a stage; i : stage number
- Has an edge from node (α, i) to two nodes $(\beta, i + 1)$ where
 - $\beta = \alpha$ by a cyclic left shift
 - $\beta = \alpha$ by a cyclic left shift + inversion of the LSBit



Example:

$(010, 1)$

→ $(100, 2)$ and

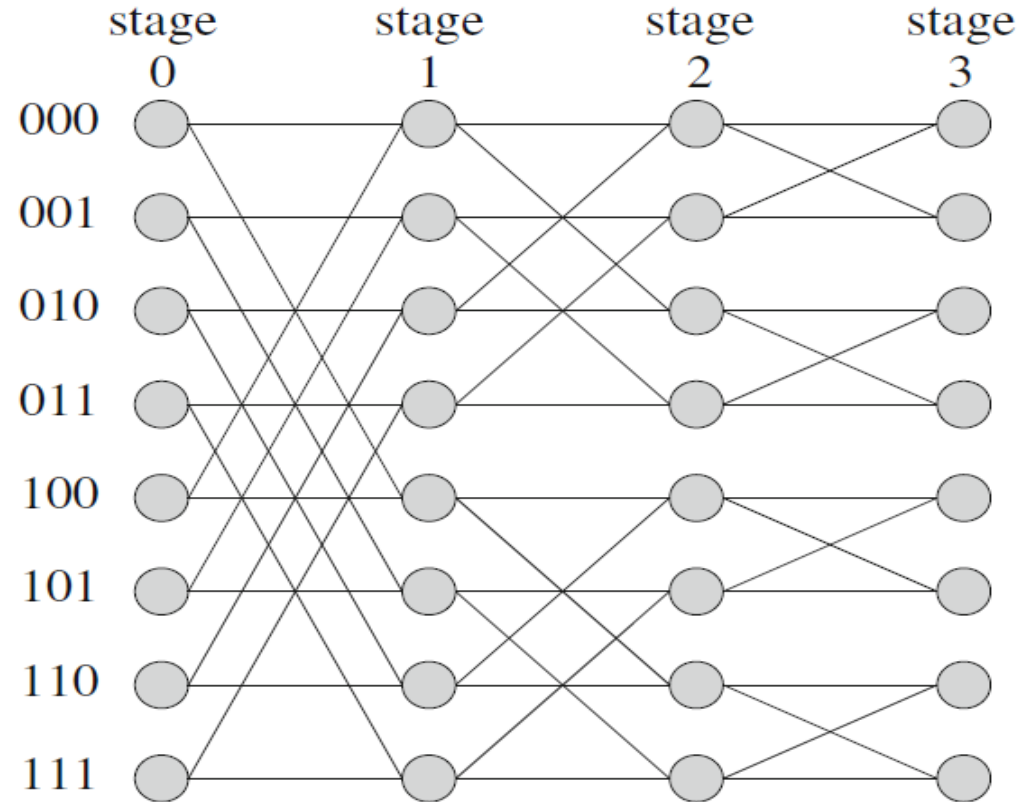
→ $(101, 2)$

**16 × 16 Omega Network
using 2x2 switches**

Omega Network vs Crossbar Switches

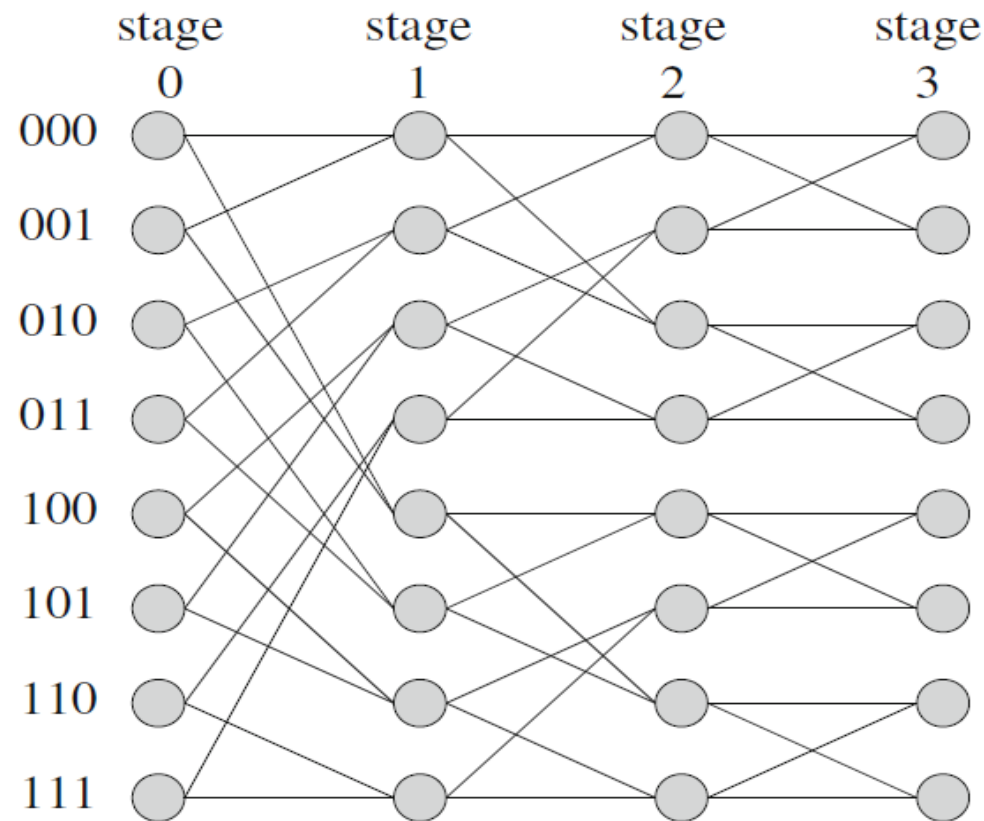
- To connect 16 processor nodes to 16 memory nodes
- Crossbar = $16 \times 16 = 256$ switches
- Omega: $n=16$ and using 2×2 switches
 - Total number of switches = $n/2$ switches per stage $\times \log n$ stages
 - 32 switches

Butterfly Network



- Node (α, i) connects to
 1. $(\alpha, i+1)$, i.e. straight edge
 2. $(\alpha', i+1)$, α and α' differ in the $(i + 1)$ th bit from the left, i.e. cross edge

Baseline Network



- Node (α, i) to **two** nodes $(\beta, i + 1)$ where
 1. β = cyclic right shift of last $(k-i)$ bits of α
 2. β = inversion of the LSB of α + cyclic right shift of last $(k-i)$ bits

Many others


- Beneš network
- Clos network
- Fat-Tree
- Folded Butterfly
- etc.....

To go from Point A to Point B

ROUTING

Routing Overview

- Routing algorithm determines path(s) from source to destination
 - Within a given interconnection topology



Routing Algorithms Classification

- Based on **path length**:


- **Minimal** or **Non-minimal** routing: whether the **shortest path** is always chosen

- Based on **adaptivity**:

- **Deterministic**: Always use the same path for the same pair of (source, destination) node
- **Adaptive**: May take into account of network status and adapt accordingly, e.g. avoid congested path, avoid dead nodes etc

- We look at three deterministic examples


XY Routing for 2D Mesh




- $(X_{\text{src}}, Y_{\text{src}})$ to $(X_{\text{dst}}, Y_{\text{dst}})$:
 - Move in X direction until $X_{\text{src}} == X_{\text{dst}}$
 - Move in Y direction until $Y_{\text{src}} == Y_{\text{dst}}$

E-Cube Routing for Hypercube

- Let $(\alpha_{n-1} \alpha_{n-2} \dots \alpha_1 \alpha_0)$ and $(\beta_{n-1} \beta_{n-2} \dots \beta_1 \beta_0)$ be the bit representations of source and destination node address respectively:
 - Number of bits difference in source and target node address → number of hops
 - Also known as *hamming distance*
 - Start from MSB to LSB (or LSB to MSB)
 - Find the first different bit
 - Go to the neighboring node with the bit corrected
- **At most n hops**




XOR-Tag Routing for Omega Network



8x8 Omega Network using 2x2 switches

- Let $T = \text{Source Id} \oplus \text{Destination Id}$
- At stage- k :
 - ❑ Go straight if bit k of T is 0
 - ❑ Crossover if bit k of T is 1



For your exploration

■ Routing:

- ❑ Adaptive Routing Algorithm
- ❑ Deadlock and deadlock avoidance in routing

■ Switching:

- ❑ Protocols
- ❑ Splitting and reconstruction of message

■ Flow Control:


- ❑ Mechanisms to avoid network congestion

So, what's the "hotness" now?

CURRENT TREND


Top 500 Supercomputers: Interconnect

Interconnect Family System Share





Interconnect Performance Share

Interconnect Family Performance Share




Ethernet

- Support *point-to-point* and *broadcast*
- Commonly-used topologies: physical bus, physical star with a logical bus, etc




InfiniBand

- Support *point-to-point* and *multicast* (on top of point-to-point capabilities)
- Commonly-use topologies: Fat tree , torus, etc.




Two level fat tree



3D torus

InfiniBand Roadmap



Summary

- Interconnection networks
 - Design criteria and properties of interconnection networks
 - Direct and indirect interconnection networks
- References:
 - Infiniband Overview:
 - <https://www.infinibandta.org/ibta-specification/>